



INF560

Algorithmique Parallèle et Distribuée

2021/2022

Patrick CARRIBAULT

CEA, DAM, DIF, F-91297 Arpajon



Lecture Outline - OpenMP

- ▶ OpenMP Basics
 - Introduction
 - Parallel region
 - Data flow
 - Worksharing
- ▶ Synchronization
 - Introduction
 - Barrier
 - Atomic operation
 - Critical region
 - Exclusive execution
 - Locks

OpenMP Basics

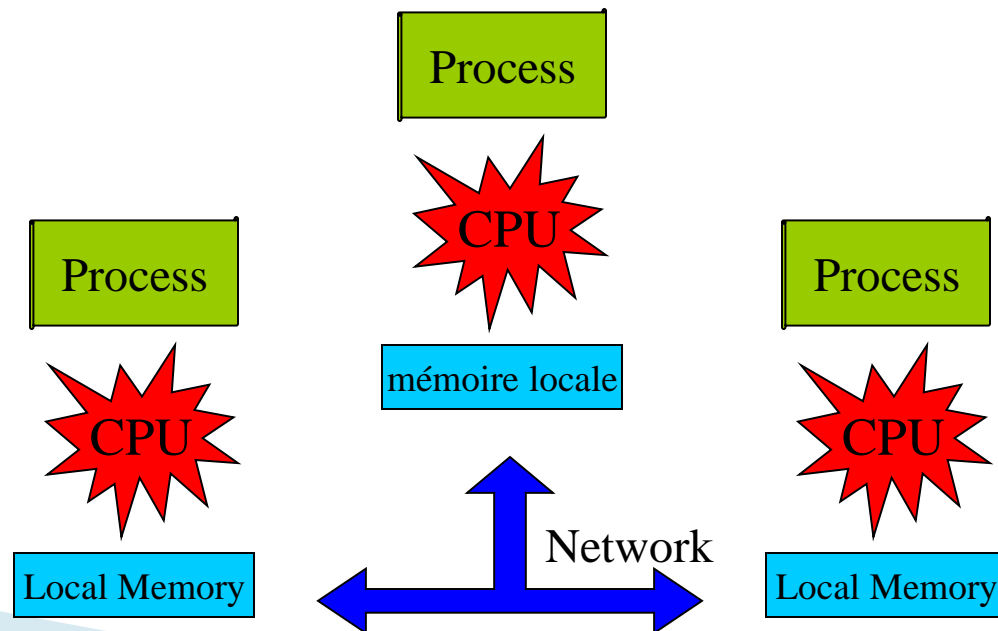
»» Introduction

Introduction

- ▶ OpenMP (Open Multi-Processing)
 - Programming model for parallel computing on shared-memory system
 - Supported on numerous platforms
 - Unix, Linux, Windows
 - Enabled on multiple programming languages
 - C/C++ and FORTRAN
 - Set of directives + library + environment variables
- ▶ Portable and scalable
 - Fast fine-grain parallelization
 - Support multiple parallelism types
 - Allow scalability (depending on the constructs used in the target application)
 - Main goal: enable parallelism for a whole node!
- ▶ Hybrid MPI+OpenMP parallelism can be used

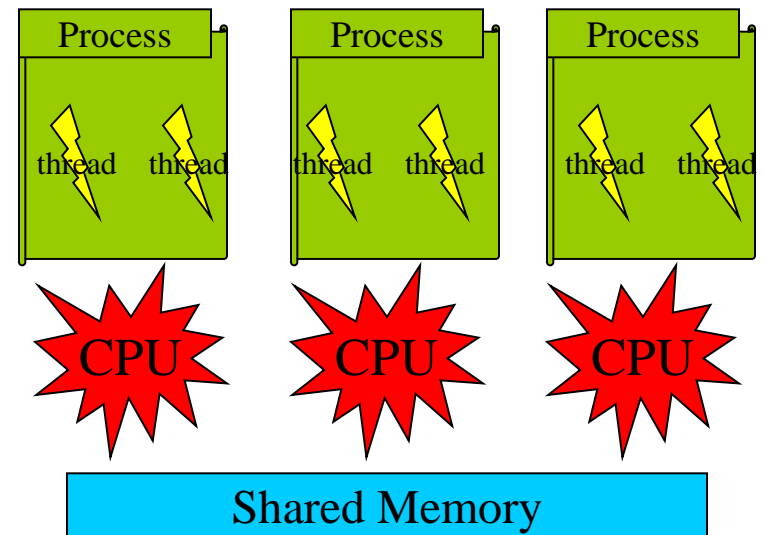
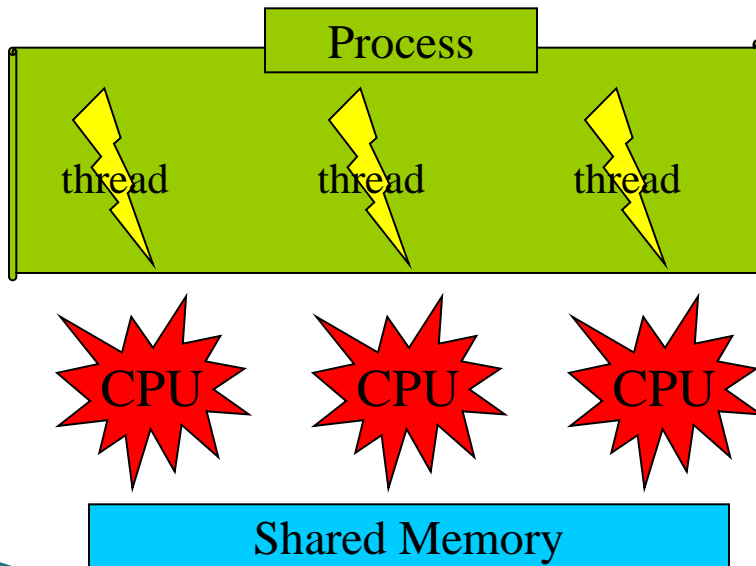
Distributed-Memory System

- ▶ Distributed-memory system
 - System in which multiple compute units have their own memory space
 - Units cannot directly access other memories
 - Network may represent inter-socket or inter-core or inter-node connections



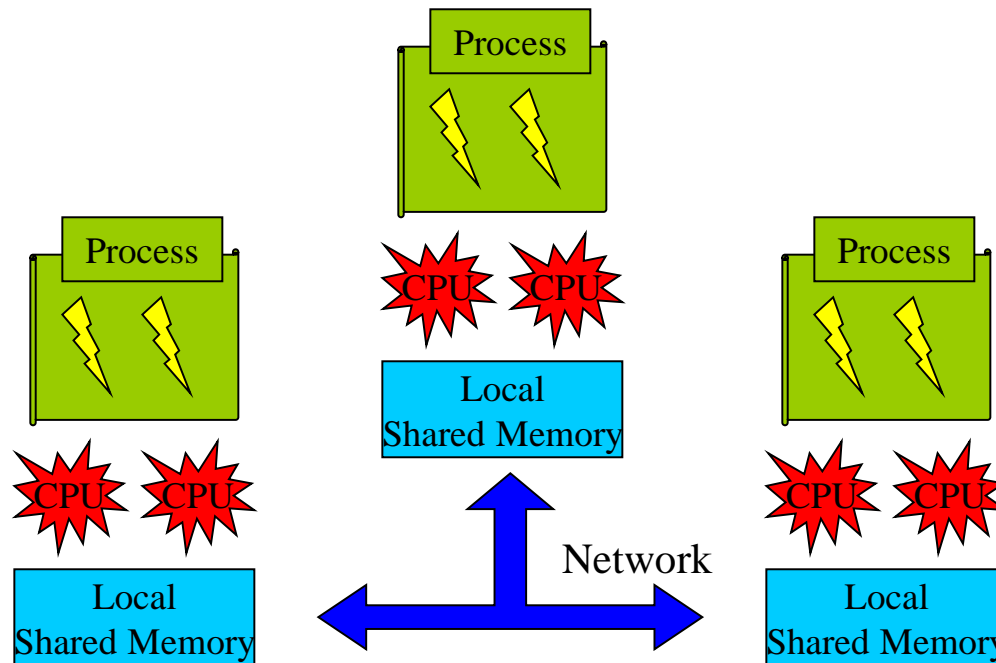
Shared-Memory System

- ▶ Shared-memory system
 - System in which multiple compute units share memory
 - Logical or physical
- ▶ Node
 - Largest set of units sharing memory



Mixed Systems

- ▶ Mix of shared and distributed memory systems
- ▶ Cluster
 - Set of nodes linked with network



Shared-Memory Model

- ▶ Requirement
 - **Parallel tasks should have the same view of memory**
- ▶ Consequence
 - Concurrent accesses to memory should be handled
- ▶ Simple approach but may lead to performance decrease
 - Critical sections/parts are sequential (by definition)
 - Data locality may not be optimal

Shared-Memory Model

- ▶ On distributed-memory system
 - Difficult
 - How to share the memory view?
 - DSM (Distributed Shared Memory)
 - May generate a large overhead
 - Depend on the number of remote accesses
- ▶ On shared-memory system
 - Easy because of shared memory
 - Inside multithreaded process
 - Every thread have access to the same memory zone
 - Usually, whole node memory

Shared-Memory Model

- ▶ API POSIX pthread
 - Standard thread management inside process
 - Suitable for MPMD approach
 - Mainly C/C++
- ▶ OpenMP
 - Compiler directives
 - Hide some management complexity (thread creation, synchronization...)
 - C/C++/FORTRAN
- ▶ TBB, Cilk+...
 - Library-based approach
 - Well integrated to C++ (template, objects...)



Lecture
Focus

History

- ▶ Multitask parallelization proposed by various vendors (e.g., CRAY, NEC, IBM...)
 - Everyone provided its own set of directives
- ▶ *Standard* definition
 - Motivated by multiprocessor machine
- ▶ Tentative w/ PCF
 - Parallel Computing Forum
 - Never adopted
- ▶ Industrial & vendor consortium → OpenMP
 - October 28, 1997
 - *De facto* standard
 - Said to be industrial standard

History – Part 1

- ▶ Managed by OpenMP ARB (*Architecture Review Board*)
 - <http://www.openmp.org>
 - <http://www.compunity.org>
- ▶ OpenMP 1.0 for FORTRAN → October 1997
- ▶ OpenMP 1.0 for C/C++ → October 1998
- ▶ OpenMP 2.0 pour Fortran → 2000
- ▶ OpenMP 2.0 pour C/C++ → 2002
- ▶ OpenMP 2.5 → May 2005
 - Unified standard for FORTRAN & C/C++
- ▶ OpenMP 3.0 → May 2008
 - Task support
- ▶ OpenMP 3.1 → July 2011
 - Taskyield construct
 - Extension of atomic operations

History – Part 2

- ▶ OpenMP 4.0 → July 2013
 - SIMD constructs
 - PROC_BIND and places
 - Device constructs
 - Task dependencies
- ▶ OpenMP 4.5 → November 2015
 - Taskloop constructs
 - Task priority
- ▶ OpenMP 5 → November 2018
 - Released SC'18 conference after multiple drafts
 - Support of OMPT (tool interface)
- ▶ OpenMP 5.1 → November 2020
 - Extend language support (C11, C++20, Fortran 2008) and include C++ attribute
 - Enhancement of environment and feedback (omp_display_env function + directive error)
 - Loop transformation construct
- ▶ OpenMP 6 → November 2023 (tentative)
 - More interaction w/ C++ standard
 - State-less threads (towards workers)

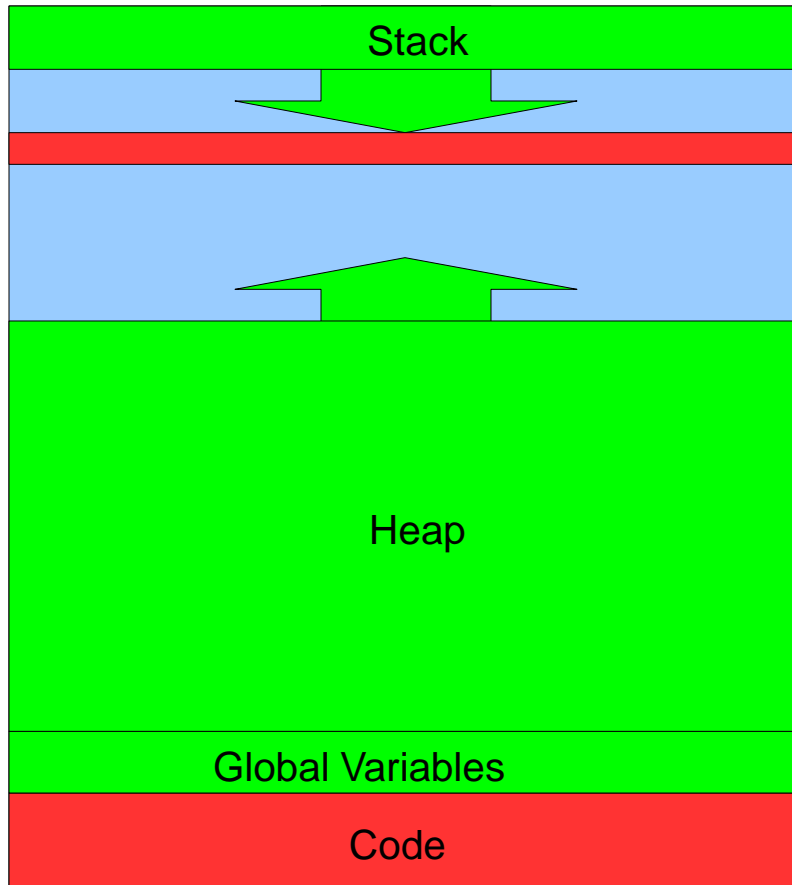
OpenMP Overview

- ▶ Programming model
 - Based on directives and API
 - Existing codes can be *augmented* with OpenMP
- ▶ Execution model
 - Execution inside one process
 - This process creates and activates threads
 - Based on the fork/join model
- ▶ Each thread
 - Executes an implicit task made of instructions
 - Has a specific rank
- ▶ Communication mode
 - Based on thread capability

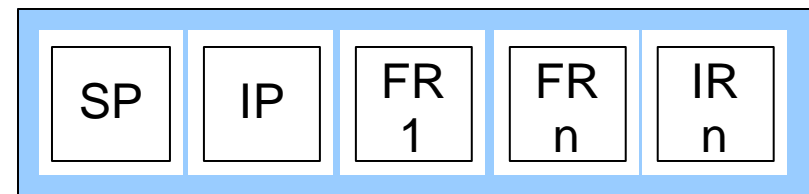
Thread vs. Process

- ▶ Thread = lightweight process
- ▶ Parts of a thread
 - Stack
 - Context
 - Include set of registers
- ▶ Parts of multi-threaded process
 - Page table
 - Set of threads

Process

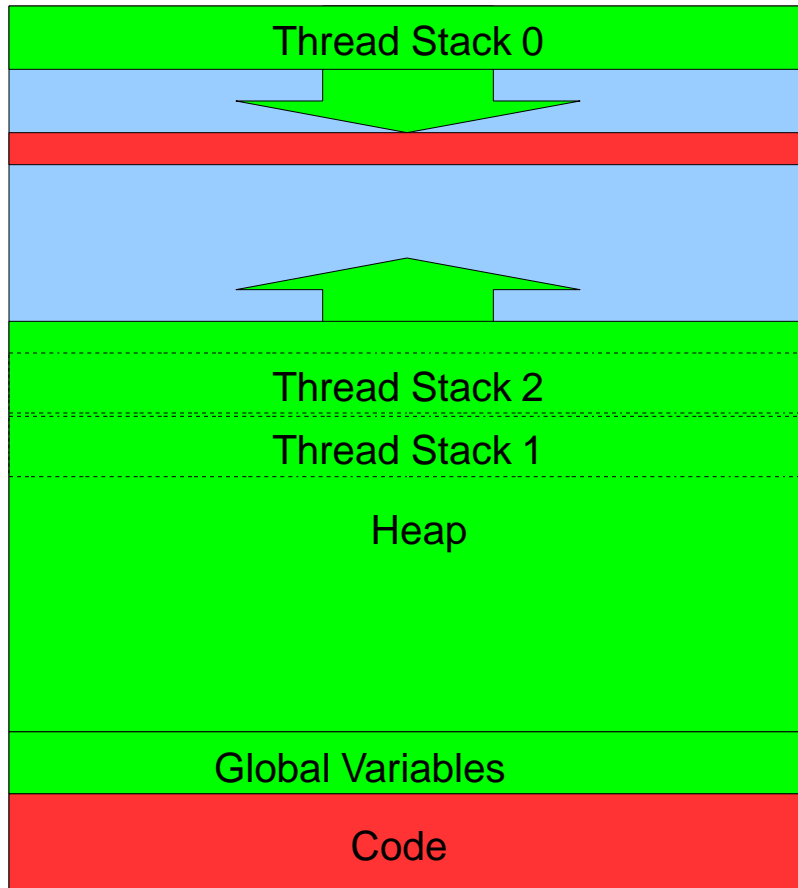


Memory

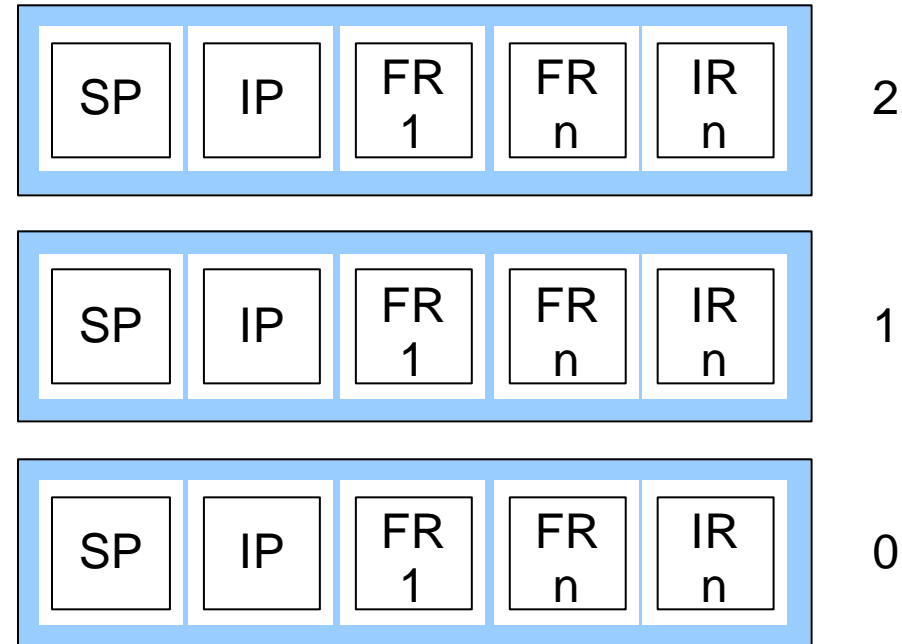


Structures

Multi-Threaded Process



Memory



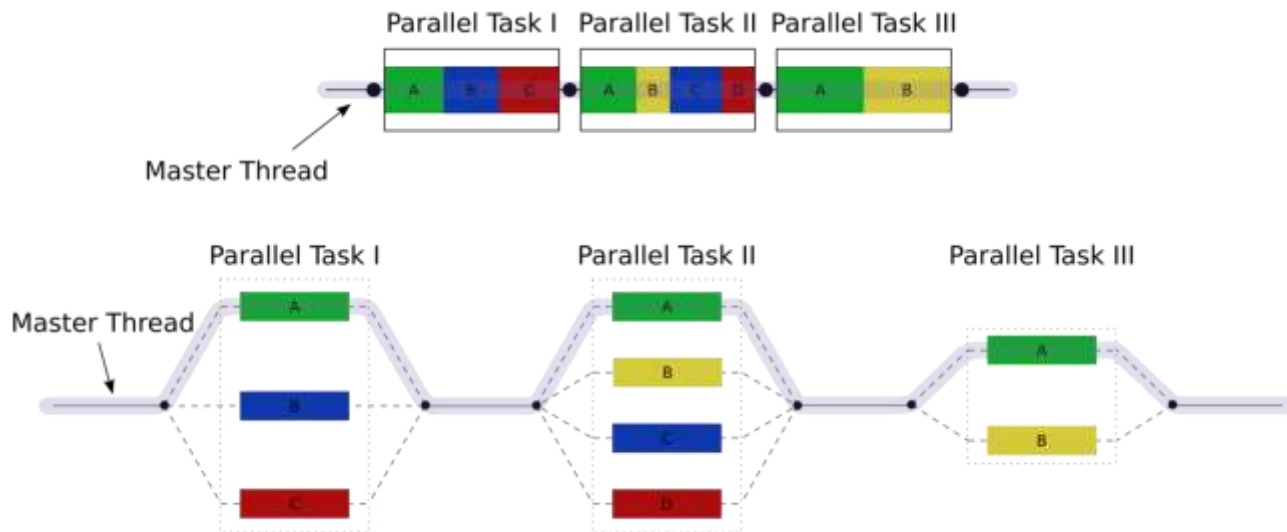
Structures

OpenMP Basics

» Parallel Region

OpenMP Execution Model

- ▶ Fork/join model
 - Program starts w/ serial part (by master thread)
 - Entering a parallel region → fork
 - Exiting a parallel region → join (barrier)
- ▶ Expressed through directives & function calls



General Syntax

- ▶ Main directive syntax for C/C++
 - `#pragma omp directive-name [clause[[,] clause] ...] new-line`
- ▶ Main directive syntax for FORTRAN
 - `sentinel directive-name [clause[[,] clause]...]`
 - Sentinel for fixed source form: `!$omp | c$omp | *$omp`
 - Sentinel for free source form: `!$omp`
- ▶ Impact of directives
 - Processed by compiler if OpenMP mode is on
 - Ignored by the compiler if OpenMP mode if off
- ▶ Library
 - Some functions are available
 - C/C++ header file: `omp.h`
 - FORTRAN module: `OMP_LIB`

Hello World!

```
#include <omp.h>  
#include <stdio.h>
```

Header

```
void main() {  
    #pragma omp parallel
```

Directive

**Parallel
Region**

```
{  
    printf( "Hello from thread %d\n",  
        omp_get_thread_num() ) ;  
}  
}
```

OpenMP functions

```
int omp_get_thread_num() ;  
int omp_get_num_threads() ;
```

Hello World!

```
void main() {  
    #pragma omp parallel  
    {  
        printf( "Hello  
                from thread %d\n",  
                omp_get_thread_num()  
                ) ;  
    }  
}
```

Execution flow

Sequential

Parallel

Sequential

Compilation

- ▶ Underlying compiler should be aware of OpenMP
- ▶ Directives are processed by the compiler
 - Need a flag to activate OpenMP support
 - Example: `-fopenmp` for GNU compiler
 - Classical bug: OpenMP compiler flag missing
- ▶ Compiler adds correct flags to link to OpenMP library
 - Link to function API
 - Link to internal functions (calls generated by compiler during directive lowering)
- ▶ More details of compilation process for OpenMP in next lecture!

Hello World!

```
#include <omp.h>
#include <stdio.h>
```

```
void main() {
    #pragma omp parallel
    {
        printf( "Hello from thread %d\n",
            omp_get_thread_num() ) ;
    }
}
```

```
$ gcc -o test -fopenmp test.c
```

```
$
```


Execution

- ▶ Code execution
 - Regular execution
 - As any multithreaded program!
- ▶ Local execution
 - Just run the executable
- ▶ Remote/Cluster execution
 - Use SLURM with one process (on one node)

Hello World!

```
#include <omp.h>
#include <stdio.h>
```

```
void main() {
    #pragma omp parallel
    {
        printf( "Hello from thread %d\n",
            omp_get_thread_num() ) ;
    }
}
```

```
$ salloc -n 1 ./test
Hello from thread 5
Hello from thread 6
Hello from thread 7
Hello from thread 4
Hello from thread 1
Hello from thread 2
Hello from thread 3
Hello from thread 0
```

Runtime Behavior

- ▶ OpenMP runtime can be controlled
 - With directives (inside program)
 - With function calls (inside program)
 - With environment variable (outside program)
- ▶ Control the number of threads
 - Environment variable `OMP_NUM_THREADS`
 - Accept an integer for the target number of threads
 - Can provide a list of integers for nested parallelism...
 - Apply to all parallel regions
 - Possibility to specify the target number of threads per parallel region
 - Clause `num_threads(int)`
 - Be careful: number of threads not guaranteed!

Hello World!

```
#include <omp.h>
#include <stdio.h>
```

```
void main() {
    #pragma omp parallel
    {
        printf( "Hello from thread %d\n",
            omp_get_thread_num() ) ;
    }
}
```

```
$ OMP_NUM_THREADS=4 salloc -n 1 ./test
Hello from thread 1
Hello from thread 2
Hello from thread 3
Hello from thread 0
```

Hello World!

```
#include <omp.h>
#include <stdio.h>
```

```
void main() {
    #pragma omp parallel num_threads(2)
    {
        printf( "Hello from thread %d\n",
            omp_get_thread_num() ) ;
    }
}
```

```
$ salloc -n 1 ./test
Hello from thread 1
Hello from thread 0
```

```
$ OMP_NUM_THREADS=4 salloc -n 1 ./test
Hello from thread 1
Hello from thread 0
```

Parallel Region Restrictions

- ▶ Before parallel region
 - Only master thread exists
 - Other threads might be asleep, but they cannot be controlled
 - No restriction on ordering when starting a parallel region
- ▶ Inside a parallel region
 - All threads execute the same block of instructions
 - Of course, dynamic control flow can be controlled with thread rank
- ▶ Implicit synchronization (barrier) at the end of the parallel region
- ▶ Branching inwards or outwards of a parallel region is forbidden
 - Example: no `goto` instruction inside a parallel region branching after the parallel region for error message processing
 - Block of parallel region should be inside the same function

Second OpenMP Program

```
#include <stdio.h>
#include <omp.h>

int main() {
    float a;
    int    p;
    a = 91680. ; p = 0;
    #pragma omp parallel
    {
        #ifdef _OPENMP
            p=omp_in_parallel();
        #endif
        printf("a: %f ; p: %d\n",a,p);
    }
    return 0;
}
```

OpenMP macro
Defined if OpenMP is processed
(value is supported version in
format YYYYMM)
→ 201511 = OpenMP 4.5
(see History section for dates and
versions)

OpenMP function
`int omp_in_parallel() ;`

```
$ export OMP_NUM_THREADS=4
$ srun -n 1 ./prog
a: 91680. ; p: 1
a: 91680. ; p: 1
a: 91680. ; p: 1
a: 91680. ; p: 1
```

OpenMP Basics

» Data Flow

Data Flow

- ▶ Data accessed inside parallel region must have a deterministic behavior
 - Depends on data scope (including declaration location)
 - Depends on data clauses
- ▶ By default → data are shared inside a parallel region
 - One variable declared before the parallel region and used inside will be shared by all threads of the region by default
 - Equivalent to clause `shared` for these variables
 - Example: `shared(a,b)`
- ▶ Default behavior can be changed with `default` clause
- ▶ Concurrent accesses to shared data
 - Be careful to concurrent read/write
 - Be careful to compiler code generation → volatility

Private Data

```
#include <stdio.h>
#include <omp.h>

int main() {
    float a;
    a = 91000.;
    printf( "Out region: %p\n", &a);
    #pragma omp parallel
    {
        printf("In region: %p thread %d\n",
            &a, omp_get_thread_num() );
    }
    return 0;
}
```

```
$ gcc -fopenmp -o test test.c
```

```
$ OMP_NUM_THREADS=4
```

```
$ salloc -n 1 ./test
```

```
Out region: 0xbf8d9a9c
```

```
In region: 0xbf8d9a9c thread 1
```

```
In region: 0xbf8d9a9c thread 2
```

```
In region: 0xbf8d9a9c thread 3
```

```
In region: 0xbf8d9a9c thread 0
```

Private Data

```
#include <stdio.h>
#include <omp.h>

int main() {
    float a;
    a = 91000.;
    printf( "Out region: %p\n", &a);
    #pragma omp parallel private(a)
    {
        printf("In region: %p thread %d\n",
            &a, omp_get_thread_num() );
    }
    return 0;
}
```

Clause

Address after
parallel region?

```
$ gcc -fopenmp -o test test.c
$ OMP_NUM_THREADS=4
$ salloc -n 1 ./test
Out region: 0xbf887e2c
In region: 0xbf887dfc thread 0
In region: 0xb67cf2ec thread 3
In region: 0xb6fd02ec thread 2
In region: 0xb77d12ec thread 1
```

Initialized Private Data

- ▶ Private data are not initialized
 - Value can be anything
- ▶ Clause to initialize data with value prior to parallel region
 - `firstprivate(var1[, var2]*)`
- ▶ Same variable behavior
 - Addresses of variables are not influenced, only values

Initialized Private Data

```
#include <stdio.h>

int main() {
    float a;
    a = 91000.;
#pragma omp parallel default(none) \
    firstprivate(a)
    {
        a = a + 680.;
        printf("a = %f\n", a);
    }
    printf("After region, a = %f\n",
        a);
    return 0;
}
```

```
$ gcc -o prog -fopenmp prog.c
```

```
$ export OMP_NUM_THREADS=4
```

```
$ salloc -n 1 ./prog
```

```
a = 91680.
```

```
a = 91680.
```

```
a = 91680.
```

```
a = 91680.
```

```
After region, a = 91000.
```

Region Extent

```
/* File prog.c */
#include <omp.h>

void sub(void);

int main() {
#pragma omp parallel
{
    sub();
}
return 0;
}
```

```
/* File sub.c */
#include <stdio.h>
#include <omp.h>

void sub(void) {
    int p=0;
#ifdef _OPENMP
    p = omp_in_parallel();
#endif
    printf("Parallel? d\n",
        p);
}
```

```
$ gcc -o prog -fopenmp prog.c sub.c
$ export OMP_NUM_THREADS=4
$ salloc -n 1 ./prog
Parallel? 1
Parallel? 1
Parallel? 1
Parallel? 1
```

Region Extent

- ▶ Local variables are private even inside other functions

```
int main() {  
#pragma omp parallel \  
    default(shared)  
{  
    sub();  
}  
return 0;  
}
```

```
#include <stdio.h>  
#include <omp.h>  
  
void sub(void) {  
    int a;  
  
    a = 91680;  
    a = a +  
        omp_get_thread_num();  
    printf("a = %d\n", a);  
}
```

```
$ salloc -n 1 ./prog  
a = 91683  
a = 91681  
a = 91682  
a = 91680
```

Arguments

- ▶ Transfer by reference acts like the original variable

- Need to check if memory space is shared or not

```
#include <stdio.h>
int main() {
    int a, b;
    a = 91000;
    #pragma omp parallel shared(a) \
        private(b)
    {
        sub(a, &b);
        printf("b = %d\n", b);
    }
    return 0;
}
```

```
$ salloc -n 1 ./prog
b = 91003
b = 91001
b = 91002
b = 91000
```

```
#include <omp.h>
```

```
void sub(int x, int *y)
{
    *y = x +
    omp_get_thread_num();
}
```


Static Variables

- ▶ Static variables
 - Global variables
 - Shared by every thread in the same process
- ▶ FORTRAN
 - Variables in COMMON block
 - Inside a MODULE
 - With SAVE qualifier
- ▶ C
 - Global variables (declared outside function blocks)
 - With static qualifier

Static Variables

```
#include <omp.h>
float a;
```

```
int main() {
    a = 91000;
#pragma omp parallel
    {
        sub();
    }
    return 0;
}
```

```
#include <stdio.h>
extern float a;
```

```
void sub(void) {
    float b;

    b = a + 680.;
    printf(
        "b = %f\n", b);
}
```

```
$ export OMP_NUM_THREADS=2
$ salloc -n 1 ./prog
b = 91680
b = 91680
```

Dynamic Allocation

- ▶ Dynamic memory allocation is allowed inside parallel region
 - Deallocation is authorized as well
- ▶ On shared variable
 - Be careful to concurrent call to dynamic allocation functions
- ▶ On private variable
 - Pointer will be initialized to a private memory zone
 - If pointer value is transmitted to other threads, this zone can be concurrently accessed

Dynamic Allocation

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
int main(){
    int n, begin, end, rank, n_tasks, i;
    float *a;
    n=1024, n_tasks=4;
    a=(float *) malloc(n*n_tasks*sizeof(float));
    #pragma omp parallel default(none) \
        private(begin,end,rank,i) shared(a,n) \
        if(n > 512)
    {
        rank=omp_get_thread_num();
        begin=rank*n;
        end=(rank+1)*n;
        for (i=begin; i<end; i++)
            a[i] = 92291. + (float) i;
        printf("Rank: %d ; A[%d],...,A[%d] : %.0f,...,%.0f \n",
            rank,begin,end-1,a[begin],a[end-1]);
    }
    free(a);
    return 0;
}
```

```
$ OMP_NUM_THREADS=4 salloc -n 1 ./prog
Rank: 3 ; A[3072],...,A[4095] : 95363.,...,96386.
Rank: 0 ; A[0000],...,A[1023] : 92291.,...,93314.
Rank: 1 ; A[1024],...,A[2047] : 93315.,...,94338.
Rank: 2 ; A[2048],...,A[3071] : 94339.,...,95362.
```

OpenMP Basics

» Worksharing

Worksharing

- ▶ Launching parallel region + access to rank & number of threads
 - Enough to parallelize an application
 - But might be complicated to implement with performance
 - Example: vector addition
- ▶ OpenMP proposes worksharing directives
 - Distribute loop iterations: for
 - Distribute blocks of instructions: sections
- ▶ Worksharing implies strong constraints
 - Work that can be shared should be independent
 - Compiler and library will not check if directives are correct!
 - Worksharing constructs should be encountered by every thread of the parallel region or none
 - Same behavior as MPI collective communication (for `MPI_COMM_WORLD`)

Parallel Loop

- ▶ Distribute iteration domain of loop nest over active threads
 - Directive `for`
 - Apply on perfect loop nest
 - Need a regular for loop (no irregular loops or while loops)
 - Useful inside a parallel region (directive `for` does not launch threads)
- ▶ Iteration scheduling can be specified
 - Clause `schedule`
 - Default choice is implementation dependent!
- ▶ Scheduling policy allows better load balancing
 - May increase overhead
 - Depend on implementation
- ▶ Synchronization
 - Barrier at the end of the loop (default)
 - Can be removed with `nowait` clause

Parallel Loop

```
#include <stdio.h>

int main( int argc, char ** argv ) {
    int N ;

    N = 10 ;

    #pragma omp parallel
    {
        int i ;
        #pragma omp for schedule(static)
        for ( i = 0 ; i < N ; i++ ) {
            printf( "Thread %d running iteration %d\n",
                    omp_get_thread_num(), i ) ;
        }
    }
    return 0 ;
}
```

```
$ gcc -fopenmp -o prog prog.c
```

```
$ salloc -n 1 ./prog
```

```
Thread 0 running iteration 0
```

```
Thread 0 running iteration 1
```

```
Thread 0 running iteration 2
```

```
Thread 3 running iteration 9
```

```
Thread 2 running iteration 6
```

```
Thread 2 running iteration 7
```

```
Thread 2 running iteration 8
```

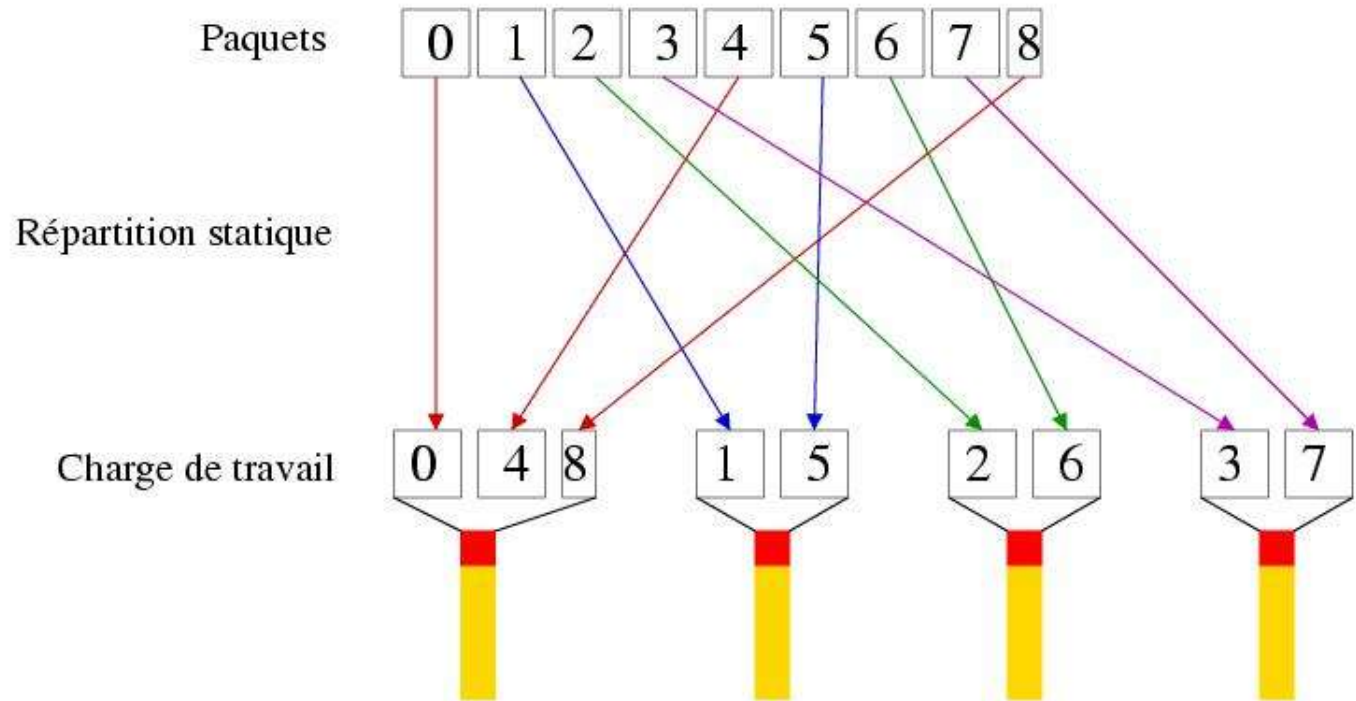
```
Thread 1 running iteration 3
```

```
Thread 1 running iteration 4
```

```
Thread 1 running iteration 5
```


Static Scheduling

- ▶ Static scheduling splits iteration domain into chunk with equal size (when possible)
 - By default: one chunk of max size
 - When specified, chunk size is constant and chunks are distributed in round robin fashion



Static Scheduling

```
#include <stdio.h>

int main( int argc, char ** argv ) {
    int N ;

    N = 10 ;

    #pragma omp parallel
    {
        int i ;
        #pragma omp for schedule(static,1)
        for ( i = 0 ; i < N ; i++ ) {
            printf( "Thread %d on iteration %d\n",
                omp_get_thread_num(), i ) ;
        }
    }
    return 0 ;
}
```

```
$ gcc -fopenmp -o prog prog.c
```

```
$ salloc -n 1 ./prog
```

```
Thread 0 on iteration 0
```

```
Thread 0 on iteration 4
```

```
Thread 0 on iteration 8
```

```
Thread 3 on iteration 3
```

```
Thread 3 on iteration 7
```

```
Thread 2 on iteration 2
```

```
Thread 2 on iteration 6
```

```
Thread 1 on iteration 1
```

```
Thread 1 on iteration 5
```

```
Thread 1 on iteration 9
```

Static Scheduling

```
#include <stdio.h>

int main( int argc, char ** argv ) {
    int N ;

    N = 10 ;

    #pragma omp parallel
    {
        int i ;
        #pragma omp for schedule(static,2)
        for ( i = 0 ; i < N ; i++ ) {
            printf( "Thread %d on iteration %d\n",
                omp_get_thread_num(), i ) ;
        }
    }
    return 0 ;
}
```

```
$ gcc -fopenmp -o prog prog.c
```

```
$ salloc -n 1 ./prog
```

```
Thread 0 on iteration 0
```

```
Thread 0 on iteration 1
```

```
Thread 0 on iteration 8
```

```
Thread 0 on iteration 9
```

```
Thread 3 on iteration 6
```

```
Thread 3 on iteration 7
```

```
Thread 1 on iteration 2
```

```
Thread 1 on iteration 3
```

```
Thread 2 on iteration 4
```

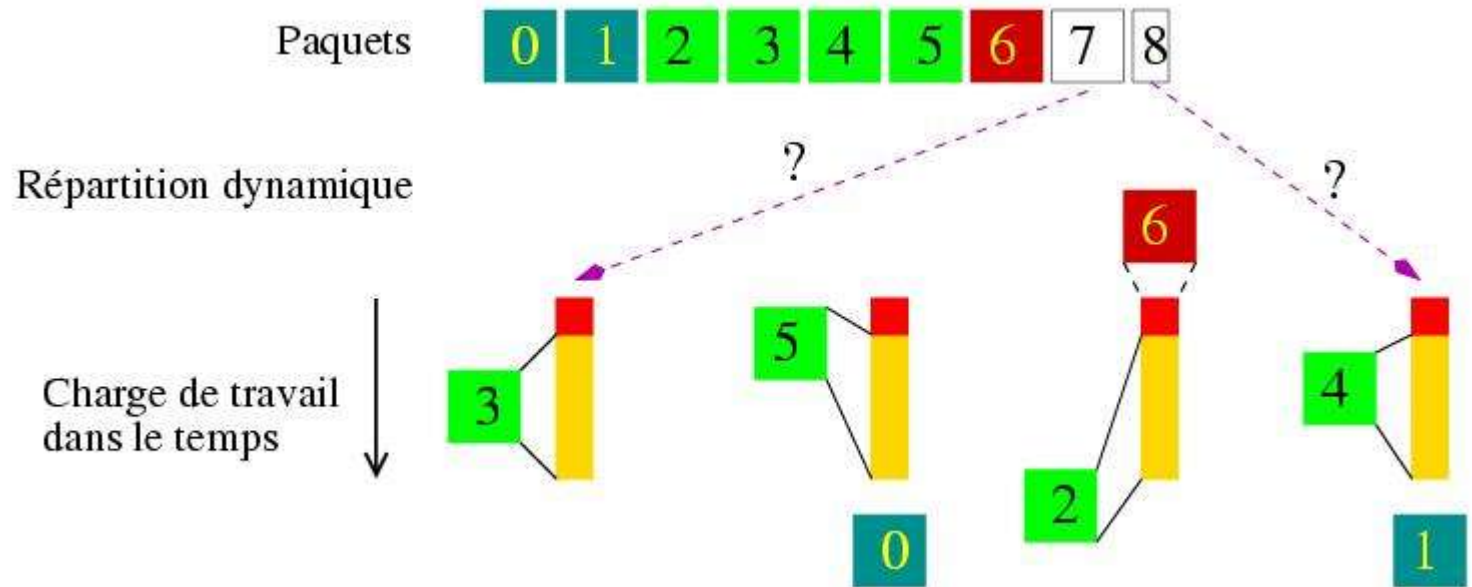
```
Thread 2 on iteration 5
```

Schedule Clause

- ▶ Scheduling policy can be controlled at runtime
 - Need to specify `schedule(runtime)` for target loops
 - Environment variable `OMP_SCHEDULE`
 - Function call `omp_set_schedule(...)`
- ▶ Scheduling policy and chunk size influence performance
 - Depends on the loop
 - Iteration domain
 - Body size
 - Depends on the target architecture

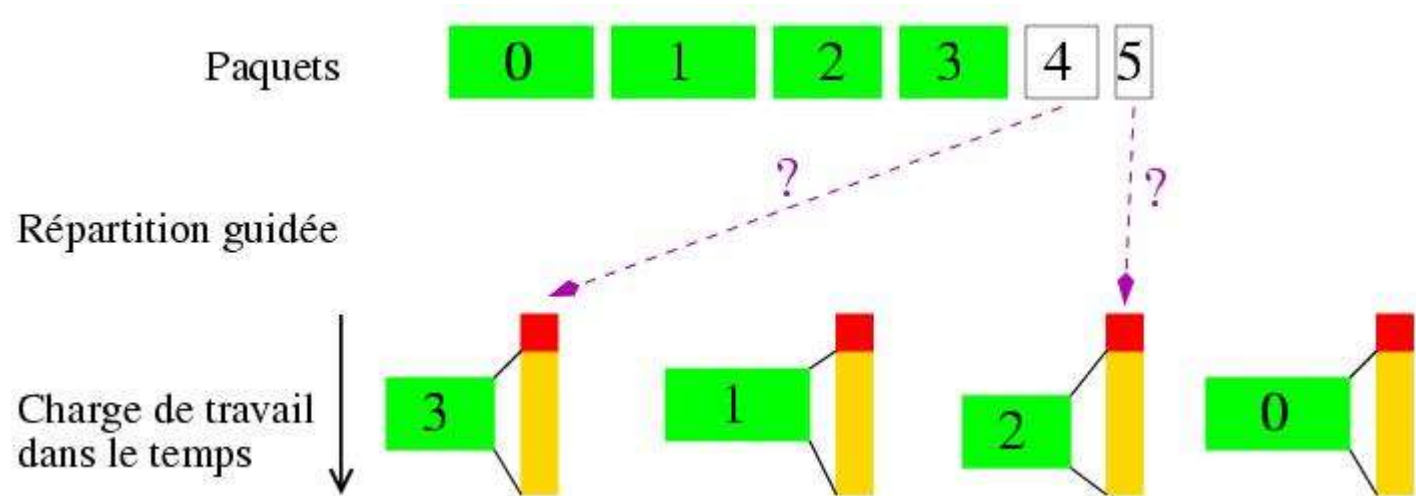
Dynamic Scheduling

- ▶ `schedule(dynamic)`
- ▶ Iterations are divided into chunks
 - Default chunk size is 1
- ▶ Chunks are distributed on-demand to threads
 - As-if behavior



Guided Clause

- ▶ Dynamic scheduling may have a large overhead (because of chunk size)
 - Large chunk → less scheduling overhead but less load balancing
 - Small chunk → Load balancing but large scheduling overhead
- ▶ Possible solution
 - Schedule(guided)
- ▶ Iterations are packed into chunks with variable size
 - Start with large chunks
 - Finish will small chunks (minimal size can be specified)



Reduction

- ▶ OpenMP Reduction
 - Associative operation applied on shared variable
- ▶ Available operations
 - Arithmetic : +, -, * ;
 - Logic : .AND., .OR., .EQV., .NEQV. ;
 - Function : MAX, MIN, IAND, IOR, IEOR.
- ▶ Each thread automatically computes partial result
 - Global result is computed with implementation-dependent strategy
 - GCC: sequential aggregation
 - INTEL: tree-based aggregation

Reduction

```
#include <stdio.h>
#define N 5
int main()
{
    int i, s=0, p=1, r=1;

    #pragma omp parallel
    {
        #pragma omp for reduction(+:s) reduction(*:p,r)
        for (i=0; i<N; i++) {
            s = s + 1;
            p = p * 2;
            r = r * 3;
        }
    }
    printf("s = %d ; p = %d ; r = %d\n",s,p,r);
    return 0;
}
```

```
$ gcc -o prog -fopenmp prog.c
```

```
$ export OMP_NUM_THREADS=4
$ salloc -n 1 ./prog
s = 5 ; p = 32 ; r = 243
```


Data Flow in Loops

- ▶ Loop construct allows clauses for data flow
- ▶ Private
 - Scoping: local declaration
 - Initial value: unknown
- ▶ Firstprivate
 - Scoping: same as private
 - Initial value: value before the loop
- ▶ Lastprivate
 - Scoping: same as private
 - Initial value: value after last iteration of the loop

Combined Regions

- ▶ Some directives may be merged together
- ▶ Examples:
 - `#pragma omp parallel`
 - `#pragma omp for`
 - ➔
 - `#pragma omp parallel for`
- ▶ Clauses
 - Union of possible clauses for both directives
- ▶ End of directive implies global synchronization (end of parallel region)

Loop Nest

- ▶ Loop directive may apply on loop nest
 - Clause `collapse(int)`
 - Restrictions: perfect loop nest
- ▶ Argument
 - Nest depth (i.e., number of loops that will be distributed over threads)
 - Available since OpenMP 3.0

Synchronization

»» Introduction

Synchronizations

- ▶ Threads launched by a parallel region have no constraints
 - Threads start at the beginning of parallel region
 - They stop at the end of the parallel region

- ▶ But synchronization mechanisms might be used
 1. To ensure control checkpointing (global barrier)
 2. To be sure that some pieces of code are executed by only one thread at a time (mutual exclusion).
 3. To synchronize two tasks (locks).

Synchronization

» Barrier

Barrier

- ▶ Synchronization of control flow between all threads inside the same team
- ▶ Implicit barrier
 - (Almost) every construct involve an implicit barrier at the end
 - Synchronize all threads from the same team
- ▶ Explicit barrier
 - Directive `barrier`
`#pragma omp barrier`

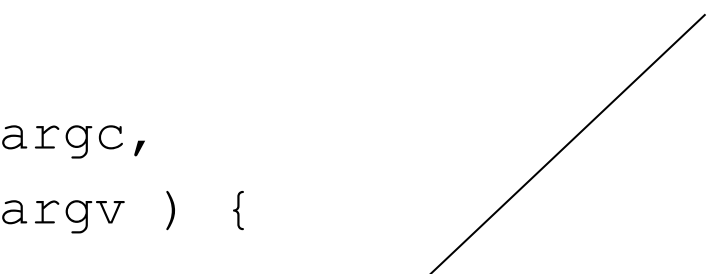
Synchronization

» Atomic Operations

Atomicity

```
int i = 0 ;

int main(int argc,
        char ** argv ) {
    i += argc ;
    printf( "%d\n", i ) ;
}
```



```
movl %esp, %ebp
movl i, %eax
addl 8(%ebp), %eax
```

Atomicity in OpenMP

- ▶ Construct for atomicity
- ▶ Ensure that following statement is executed atomically

```
#pragma omp atomic
```

- ▶ Following statement can be of specified form (usually variable update with arithmetic operation)

Atomicity in OpenMP

```
#include <stdio.h>
#include <omp.h>

int main() {
    int c, r;

    c = 91680;
    #pragma omp parallel private(r)
    {
        r = omp_get_thread_num();

        #pragma omp atomic
        c++;

        printf("R: %d ; C: %d\n", r, c );
    }
    printf("After region, c: %d\n", c);
    return 0;
}
```

```
$ gcc -o prog -fopenmp prog.c
```

```
$ export OMP_NUM_THREADS=4
```

```
$ salloc -n 1 ./prog
```

```
R: 1 ; C: 91683
```

```
R: 0 ; C: 91681
```

```
R: 2 ; C: 91684
```

```
R: 3 ; C: 91682
```

```
After region, c: 91684
```

Atomicity in OpenMP

- ▶ Atomic statement

- $x = x \sim (\text{op}) \sim \text{exp}$;
- $x = \text{exp} \sim (\text{op}) \sim x$;
- $x = f(x, \text{exp})$;
- $x = f(\text{exp}, x)$;

- ▶ (op) should be

- $+$, $--$, $*$, $/$,
- .AND. , .OR. , .EQV. , .NEQV. .

- ▶ f is a function

- MAX , MIN , IAND , IOR , IEOR .

- ▶ exp expression that does not depend on x

Synchronization

» Critical Region

Critical region

- ▶ Extension of atomic operation to a region (i.e., block of statements)
 - Threads execute the inner block one at a time
 - Dynamic extent
- ▶ Directive followed by a block

```
#pragma omp critical
```
- ▶ Possibility to create named critical region
 - By default, region is anonymous
 - It applies to every region with the same name (including anonymous)
- ▶ Atomic operation and critical region have different performance...

Critical region

```
#include <stdio.h>

int main()
{
    int s, p;

    s = 0, p = 1;
    #pragma omp parallel
    {
        #pragma omp critical
        {
            s++;
            p*=2;
        }
    }
    printf("Sum & product: %d, %d\n", s, p);
    return 0;
}
```

Synchronization

» Exclusive Execution

Exclusive Execution

- ▶ Mechanism to allow only one thread to execute a target block
 - Block executed only once
- ▶ OpenMP proposes 2 directives
 - `#pragma omp master`
 - `#pragma omp single`
- ▶ Main goal is the same but definition, code generation and runtime behavior are different...

Master Construct

- ▶ Define a block of instructions that will be executed only by master thread
- ▶ No clauses
 - Directive is followed by a block of statements
- ▶ No synchronization at the entrance or at the end of the master block

Master Construct

```
#include <stdio.h>
#include <omp.h>

int main()
{
    int r;
    float a;

    #pragma omp parallel private(a,r)
    {
        a = 91680.;

        #pragma omp master
        {
            a = -91680.;
        }

        r = omp_get_thread_num();
        printf("R: %d ; A: %f\n",r,a);
    }
    return 0;
}
```

Single construct

- ▶ Define a block of instructions to be executed only by one thread
 - Generalized master
 - `#pragma omp single`
- ▶ No constraints on which thread will execute the block
 - Maybe the first one
- ▶ Implicit barrier at the end of single construct
 - Single construct can be seen as collective operation!
 - Clause `nowait` to remove this barrier

Single construct

```
#include <stdio.h>
#include <omp.h>

int main()
{
    int rang;
    float a;

    #pragma omp parallel private(a,rang)
    {
        a = 91680.;

        #pragma omp single
        {
            a = -91680.;
        }

        rang=omp_get_thread_num();
        printf("Rang : %d ; A vaut : %f\n",rang,a);
    }
    return 0;
}
```

Synchronization

» Locks

Lock Definition

- ▶ Locks: object for mutual exclusion

- Allow shared-data protection
- Enable critical sections

- ▶ States

- locked (owned by a thread)
- unlocked (no ownership)

- ▶ Actions

- Acquire the lock (lock)
- Release the lock (unlock)

- ▶ State update

- Lock on locked status → wait for release
- Lock on unlocked status → update status to locked
- Unlock on unlocked status → nothing (not recommended)
- Unlock on locked status → release

2-bit FSA

The diagram consists of two blue rectangular boxes with double borders. The top box is labeled '2-bit FSA'. The bottom box is labeled 'Nested/Recursive: Allow same owner multiple times'. A blue line with a double border starts from the 'State update' section, goes up and to the right, then turns down and to the right, ending at the bottom box. Another blue line with a double border starts from the 'State update' section, goes up and to the right, then turns down and to the right, ending at the top box.

Nested/Recursive:
Allow same owner
multiple times

OpenMP Locks

▶ Opaque type for locks

- `omp_lock_t`

▶ Lock initialization

- `void omp_init_lock(omp_lock_t *lock);`

▶ Lock destruction

- `void omp_destroy_lock(omp_lock_t *lock);`

▶ Acquire a lock

- `void omp_set_lock(omp_lock_t *lock);`

▶ Release a lock

- `void omp_unset_lock(omp_lock_t *lock);`

▶ Try to acquire a lock

- `int omp_test_lock(omp_lock_t *lock);`

Initialization

Actions

OpenMP Nested Locks

- ▶ Extension to nested locks
- ▶ Opaque type for locks
 - `omp_nest_lock_t`
- ▶ Lock initialization
 - `void omp_init_nest_lock(omp_nest_lock_t *lock);`
- ▶ Lock destruction
 - `void omp_destroy_nest_lock(omp_lock_t *lock);`
- ▶ Acquire a lock
 - `void omp_set_nest_lock(omp_nest_lock_t *lock);`
- ▶ Release a lock
 - `void omp_unset_nest_lock(omp_nest_lock_t *lock);`
- ▶ Try to acquire a lock
 - `int omp_test_nest_lock(omp_nest_lock_t *lock);`

Initialization

Actions

OpenMP Lock Example

```
#include <omp.h>
#include <stdio.h>

int main()
{
    int n ;
    int c ;
    omp_lock_t l ;
    omp_init_lock( &l ) ;

#pragma omp parallel
{
    #pragma omp single
    {
        n = omp_get_num_threads() ;
    }
}
```

```
omp_set_lock(&l) ;

c++;

omp_unset_lock(&l) ;
}

omp_destroy_lock( &l ) ;

printf( "Number of threads
function:%d
count:%d\n", n, c ) ;

return 0 ;
}
```

Protect
variable c

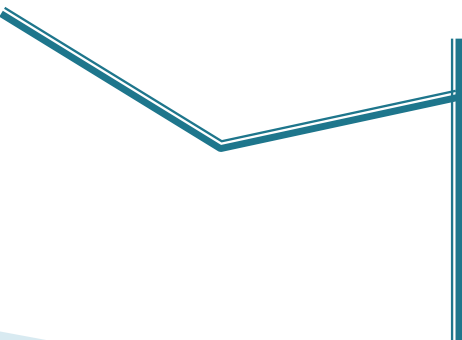
Generalized way to perform atomic
operations and critical sections

Lock Internals

- ▶ Example of lock implementation through a *mutex*
- ▶ Main structure + slot for thread queue

```
typedef struct slot_s {  
    thread_t *thread;  
    struct slot_s *next;  
} slot_t;
```

```
typedef struct {  
    /* Counter for waiting threads */  
    volatile int nb_threads;  
    /* List of blocked threads */  
    volatile slot_t *list_first;  
    volatile slot_t *list_last;  
    /* Spinlock to control accesses to internal variables */  
    spinlock_t lock;  
} mutex_t;
```



Need a lock to check
read/write accesses to
the mutex structure!

Mutex Internals

► Function to acquire the lock

```
void mutex_lock (mutex_t * m)
{
    slot_t slot;
    spinlock (&(m->lock));
    if (m->nb_thread == 0) {
        m->nb_thread = 1;
        spinunlock (&(m->lock));
    } else {
        slot.thread = thread_self ();
        enqueue (&slot, m);
        thread_self ()->status = blocked;
        register_spinunlock (&(m->lock));
        yield ();
    }
}
```

**nb_thread can be
safely checked
because of spinlock**

**Call to main scheduler
because current thread
is blocked**

Mutex Internals

► Function to release the lock

```
void mutex_unlock (mutex_t * m)
{
    slot_t *slot;
    spinlock (&(m->lock));
    if (m->list_first != NULL) {
        slot = dequeue (m);
        wake (slot->thread);
    } else {
        m->nb_thread = 0;
    }
    spinunlock (&(m->lock));
}
```

`list_first` can be
safely checked
because of spinlock

Thread releasing the
lock will wake up the
next thread waiting for
the mutex

Mutex Internals

► Function to test lock status

```
int mutex_trylock (mutex_t * m)
{
    slot_t slot;
    spinlock (&(m->lock));
    if (m->nb_thread == 0) {
        m->nb_thread = 1;
        spinunlock (&(m->lock));
        return 0;
    }
    spinunlock (&(m->lock));
    return 1;
}
```

Recursive Mutex Internals

- ▶ Function to acquire the lock

```
void mutex_lock (mutex_t * m) {
    slot_t slot;
    spinlock (&(m->lock));
    if (m->nb_thread == 0) {
        m->nb_thread = 1;
        m->owner = thread_self ();
        spinunlock (&(m->lock));
    } else {
        if (m->owner == thread_self ()) {
            m->step++;
            spinunlock (&(m->lock));
        } else {
            slot.thread = thread_self ();
            enqueue (&slot, m);
            thread_self ()->status = blocked;
            register_spinunlock (&(m->lock));
            yield ();
        }
    }
}
```

Recursive Mutex Internals

- ▶ Function to release the lock

```
void mutex_unlock (mutex_t * m) {
    slot_t *slot;
    spinlock (&(m->lock));
    if (m->step == 1) {
        m->step--;
        if (m->list_first != NULL) {
            slot = dequeue (m);
            wake (slot->thread);
        } else {
            m->nb_thread = 0;
        }
    } else {
        m->step--;
    }
    spinunlock (&(m->lock));
}
```


Recursive Mutex Internals

► Function to test lock status

```
int mutex_trylock (mutex_t * m){
    slot_t slot;
    spinlock (&(m->lock));
    if (m->nb_thread == 0) {
        m->nb_thread = 1;
        m->owner = thread_self ();
        spinunlock (&(m->lock));
        return 0;
    } else {
        if (m->owner == thread_self ()) {
            m->step++;
            spinunlock (&(m->lock));
            return 0;
        }
    }
    spinunlock (&(m->lock));
    return 1;
}
```