



INF560

Algorithmique Parallèle et Distribuée

2021/2022

Patrick CARRIBAULT

CEA, DAM, DIF, F-91297 Arpajon



Lecture Outline – GPU (1/2)

- ▶ Overview of Dedicated Hardware
- ▶ NVIDIA GPGPU architecture
- ▶ CUDA Programming
 - Host-Side
 - Device-Side

Dedicated Hardware

»» Overview

Motivations

- ▶ Homogeneous clusters
 - Building blocks: on-the-shelves CPUs
 - Good performance for majority of codes, but dedicate many transistors to micro-architecture
- ▶ Limitations of such clusters
 - Homogeneous large cores
 - Relate to limitation of general purpose processors
- ▶ One possible solution
 - → Towards dedicated hardware

Dedicated Hardware

▶ Definition

Distinct resource dedicated to specific operations

▶ Impact

- Aka *accelerators*
- Processor dedicated to a specific class of applications
 - Different goals from regular CPU
 - Enable better performance for this class
- Located on specific chip (or not)
 - Linked to CPU through some interconnect
 - Master/slave behavior

Interconnect Types

▶ Advantages

- Low cost
- Smooth integration

▶ Drawbacks

- Moderate performance (generic bus)
- Sub-optimal latency and bandwidth

▶ NVIDIA GPGPU, Intel Xe, AMD MI100

▶ Advantages

- Fast access
- Possibility to share resources w/ CPU

▶ Drawbacks

- Complex integration
- Processor dependency

▶ IBM Cell , AMD/ATI fusion architecture

System Bus

Processor Bus

Main Architecture

- ▶ Microarchitecture
 - Driven by target class of applications
 - Simplified compared to regular CPUs
 - No OoO execution
 - No branch predictor
 - ...
- ▶ Exploited parallelism
 - Fine-grain parallelism (in general)
 - Mostly following the *Single Instruction Multiple Data (SIMD)* form
- ▶ But
 - Major work is let to software stack and code developer!

Main Architecture

▶ Memory hierarchy

- Memory controlled by software/developer (*scratchpad memory*)
- Cache handled by hardware (part of microarchitecture like on regular CPU)
- Hardware connection between compute processors
- Large bandwidth between device on its global memory

▶ Summary

- Interesting characteristics for HPC
- Many aspects to handle *by hand*
- Direct influence on programming and execution model

Accelerator Programming

- ▶ Accelerator can be seen as slave device
- ▶ Main processor controls accesses to dedicated hardware
 - Main processor (CPU): *host*
 - Accelerator: *device*
- ▶ Main concepts
 - Data transfers
 - Remote code execution with dedicated language
 - Manual management of execution
 - Queries to accelerators

Existing Accelerators

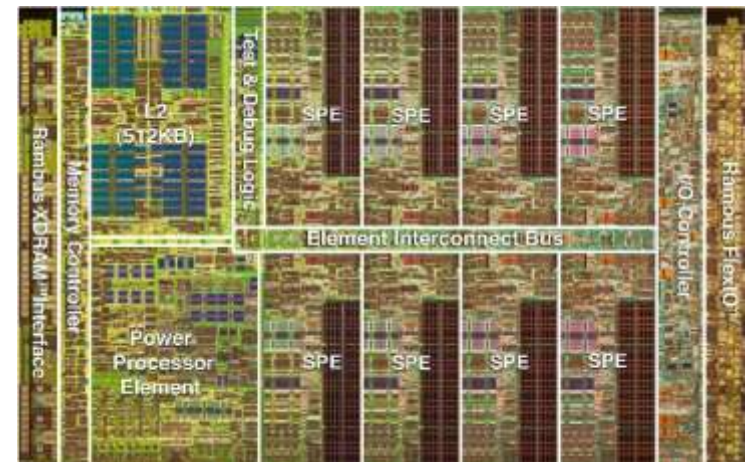
- ▶ Graphics-based processor
 - GPGPUs (General Purpose Graphics Processing Units)
 - NVIDIA
 - AMD
 - Intel Xe
- ▶ Dedicated accelerators
 - Cell (IBM)
 - Intel Xeon Phi (KNC/KNL)
- ▶ Reconfigurable architectures
 - FPGA

Graphics Processor

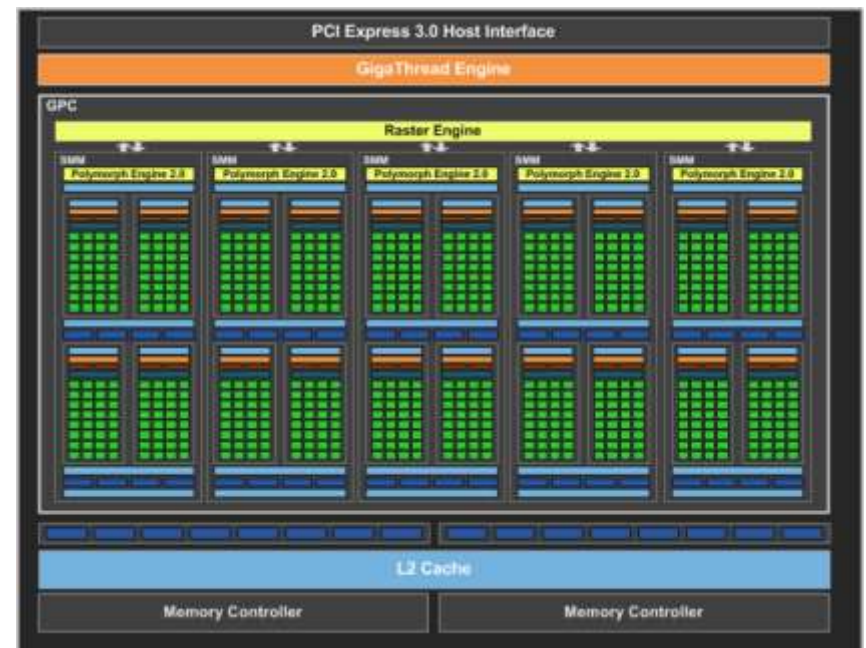

- ▶ Initially designed to print pixels on screen (for 2D or 3D graphics)
 - Dedicated processor for this purpose
- ▶ Large compute capabilities
 - Optimized to image rendering
 - Well adapted to regular data processing
- ▶ For scientific workloads
 - GPGPU (General Purpose Graphics Processing Unit)
 - Graphics cards on system bus
 - Connection to motherboard
 - Communication through PCI-E
 - Possibility to connect multiple graphics cards
 - Multi-GPU programming
 - Current configuration in heterogeneous supercomputers

IBM Cell Processor


- ▶ Joint development between IBM, Sony and Toshiba
- ▶ 8 64-bit cores (w/ FP units)
 - *Synergistic Processor Elements* (SPEs).
 - SIMD: 128-bit register width (4 words of 32 bits)
- ▶ Master processor: 64-bit PowerPC
 - 2 logical threads
- ▶ Register file
 - 128 operands




NVIDIA Maxwell Architecture



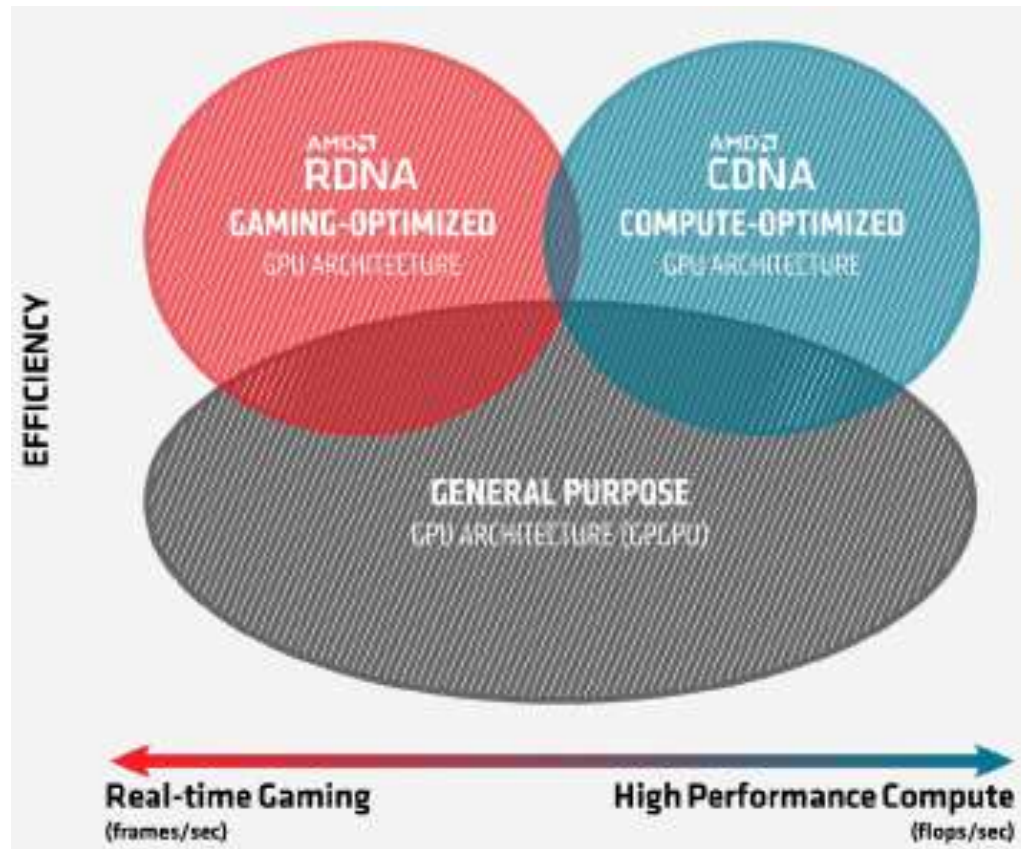
NVIDIA Volta Architecture



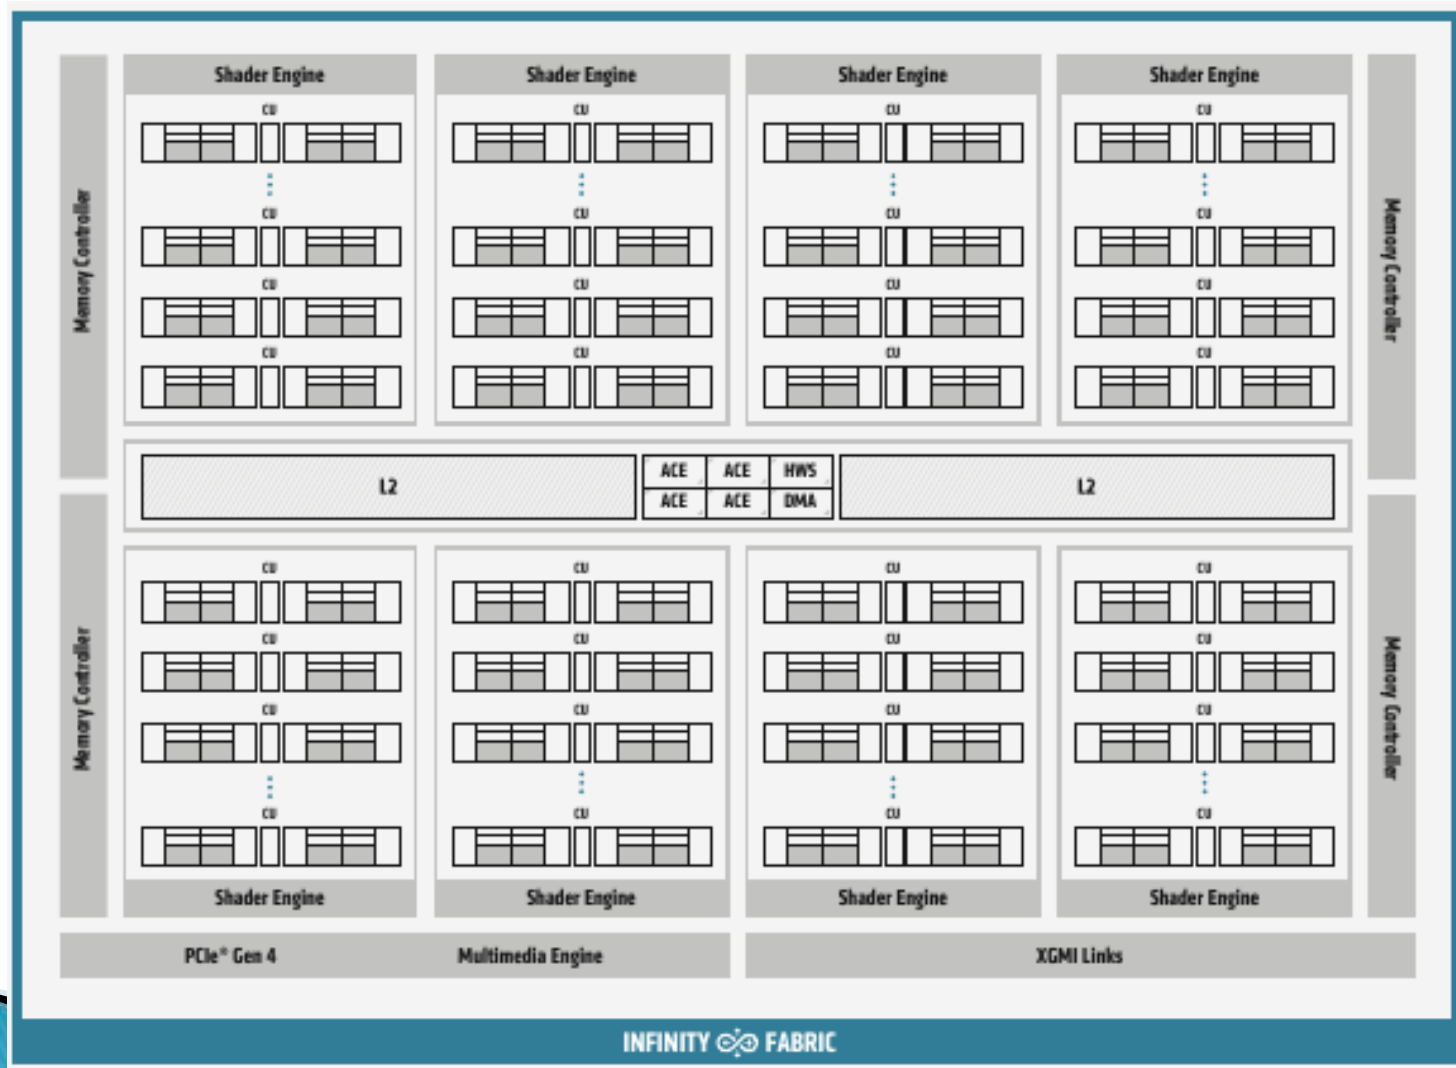
NVIDIA Ampere Architecture



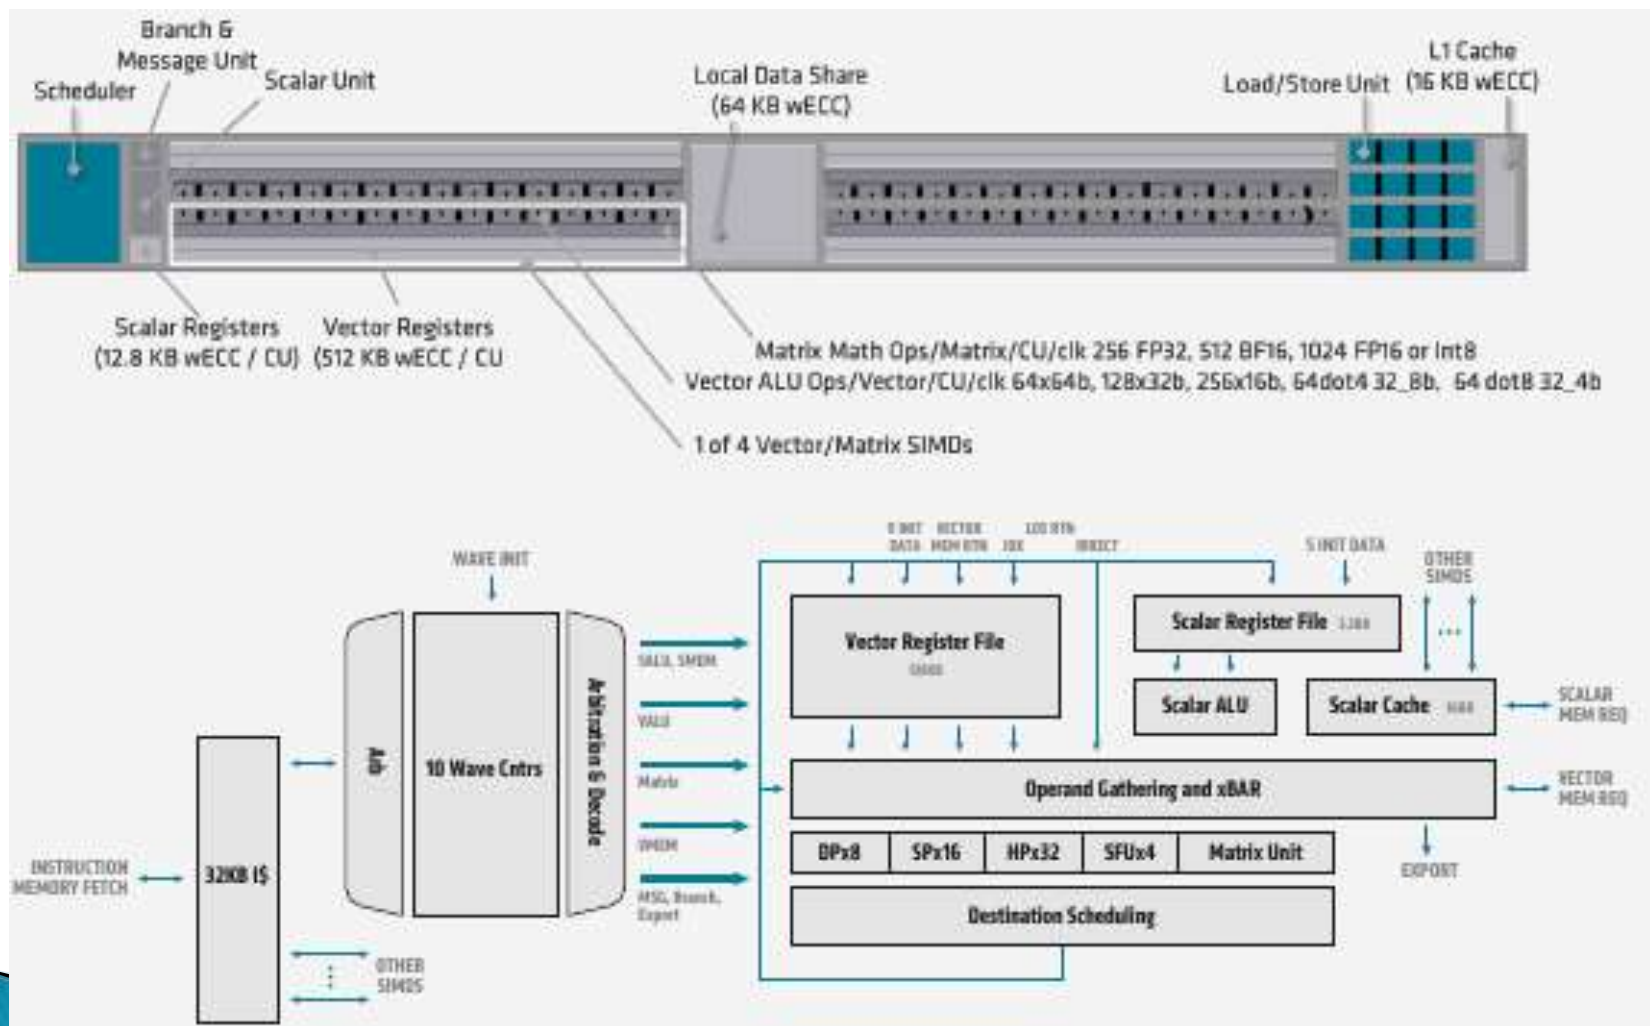
AMD CDNA Architecture



AMD CDNA Architecture




AMD CDNA Architecture



Intel Xeon Phi

- ▶ Accelerator designed by Intel
 - Many-Integrated Core (MIC)
- ▶ Introduced in 2008
 - @SigGraph 2008 conference
 - Codename at that time: *Larabee*
- ▶ Architecture
 - X86 cores
 - Support of *hyperthreading*
 - Cache coherency
- ▶ Current version: KNL
 - Exist in non-accelerator mode



Future Exascale Systems

- ▶ Next year: first exascale machine!
- ▶ Funded by DoE, located in National Labs
- ▶ Argonne National Lab: Aurora
- ▶ Oakridge National Lab: Frontier
- ▶ Based on GPUs!

ANL - Aurora

- ▶ Sustained performance > 1 Exaflop
- ▶ Aggregate system memory > 10 PB
- ▶ Storage > 230 PB, 25 TB/s (DAOS)
- ▶ Compute node
 - 2 Intel Xeon processors
 - 6 Xe arch-based GP-GPUs
- ▶ Interconnect
 - CPU-GPU: PCIe
 - System : Cray Slingshot
- ▶ Software stack
 - Cray stack (+ Intel improvement)
 - Programming models : Intel OneAPI, OpenMP, SYCL/DPC++




ANL - Aurora

Building the Foundation for Exascale Computing

2 INTEL XEON SCALABLE PROCESSORS
"Sapphire Rapids"

6 X^E ARCHITECTURE BASED GPU'S
"Ponte Vecchio"

ONEAPI
Unified programming model



LEADERSHIP PERFORMANCE
For HPC, data analytics, AI

UNIFIED MEMORY ARCHITECTURE
Across CPU & GPU

ALL-TO-ALL CONNECTIVITY WITHIN NODE
Low latency, high bandwidth

UNPARALLELED I/O SCALABILITY ACROSS NODES
8 fabric endpoints per node, DAOS

DELIVERED IN 2021

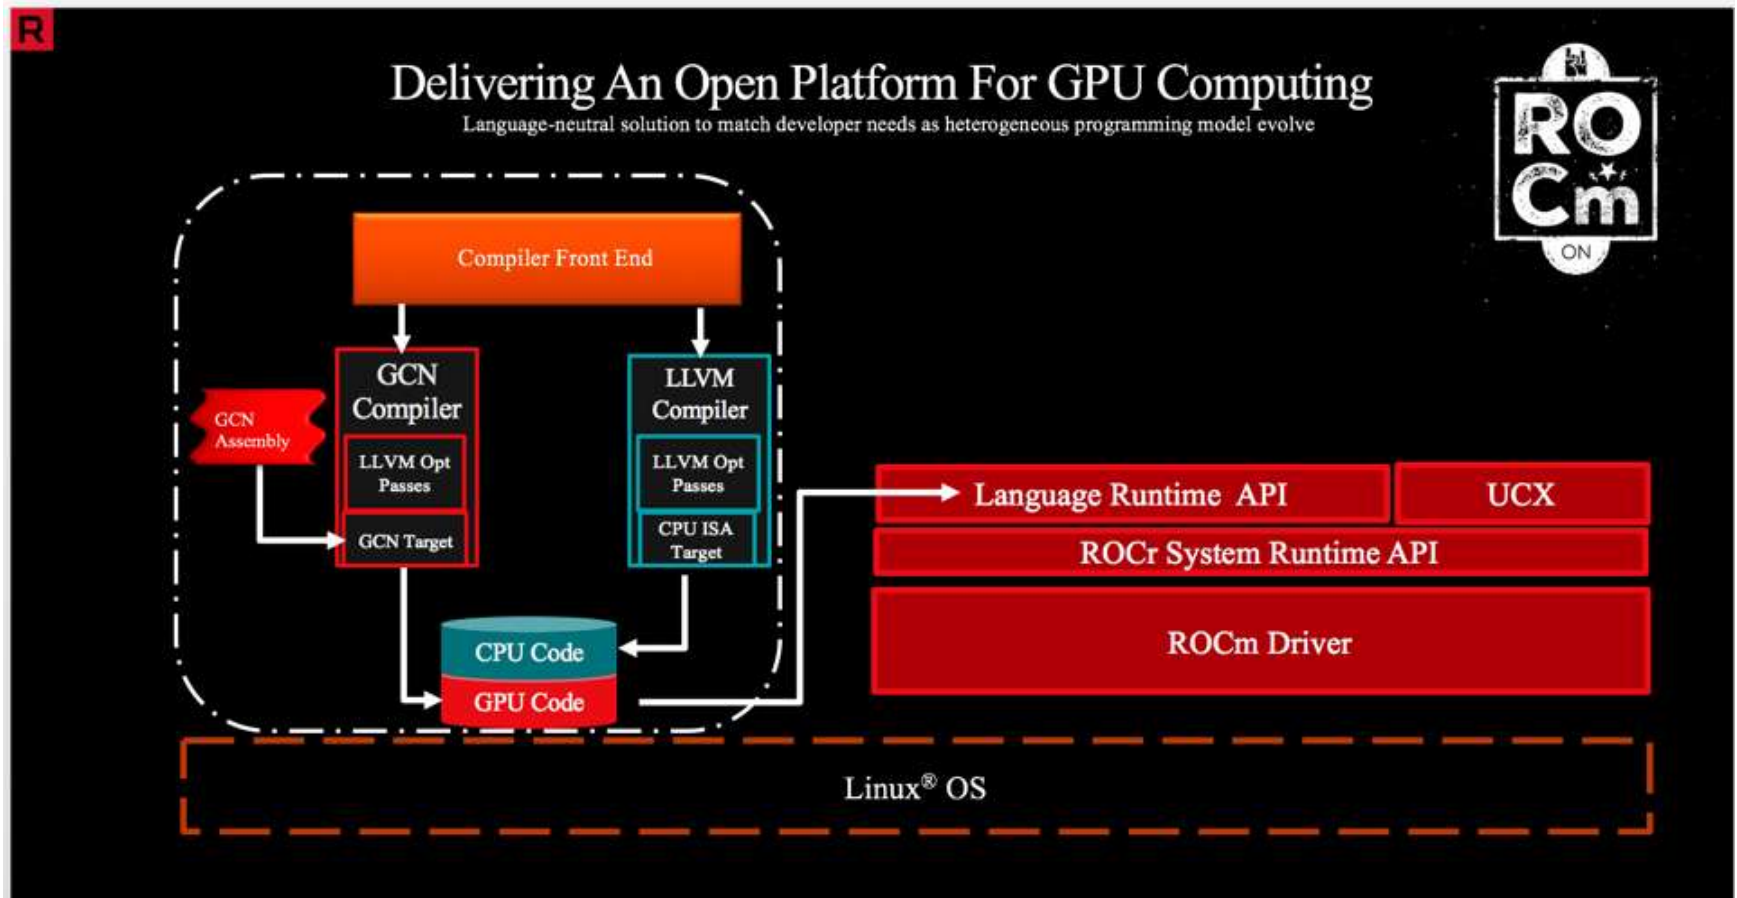


ONL - Frontier

- ▶ Peak Performance > 1.5 Exaflop
- ▶ More than 100 cabinets
- ▶ Node
 - 1 HPC and AI optimized AMD EPYC CPU
 - 4 AMD Radeon Instinct GPU
- ▶ Interconnect
 - CPU-GPUI: AMD infinity Fabrix (coherent memory across node)
 - System: Cray Slingshot
- ▶ Software stack
 - Cray + AMD ROCm
 - Programming models : MPI, OpenMP 5.X, HIP, ...




ONL - Frontier




Dedicated Hardware

»» NVIDIA GPGPU Architecture

Floating-Point Performance



Bandwidth Evolution



NVIDIA Architecture Overview

- ▶ Building block
 - Very simple compute core
 - Set of cores grouped inside a *stream multiprocessor*
- ▶ Why *stream* ?
 - Made to handle data streams (contiguous flow of data)
 - Synchronous language on each multicore processor
- ▶ Designed for fine-grain parallelism

NVIDIA Pascal Architecture

- ▶ Recent generation of NVIDIA architecture (w/ double-precision performance)
 - Codename: Pascal
- ▶ Tesla version
 - 15.3 billion transistors
 - 5.3 Tflops (double precision)



NVIDIA Pascal Architecture

- ▶ GP100 → array of 6 GPCs (Graphics Processing Clusters)
 - 1 GPC → 10 SMs (Streaming Multiprocessors)
- ▶ Placement according to L2 cache
- ▶ Figure
 - Green → compute core
 - Double-precision unit in orange
 - Orange → scheduler and dispatcher
 - Light blue → register file and L1\$



NVIDIA Pascal Architecture



Pascal: *Streaming Multiprocessor*

- ▶ 6th generation
- ▶ 64 CUDA-compatible cores
 - ALU unit
 - Simple-precision compute unit
- ▶ 32 double-precision units
- ▶ 16 special units (SFU)
- ▶ 16 load/store units




Pascal: Streaming Multiprocessor

- ▶ SIMT-based execution
 - Single-Instruction Multiple Thread
- ▶ Similar to SIMD
- ▶ Synchronous execution on compute cores
- ▶ Each SM manage threads through packs of 32
 - Notion of *warp*
 - 2 warp scheduler per SM
 - 2 dispatch units/scheduler → 2 issued instruction per cycle



Pascal: Memory Hierarchy

- ▶ Threads are scheduled on one compute core (on one SM)
- ▶ Memory access through *load/store* instruction
- ▶ 2 possibilities
 - Shared-memory access → no cache
 - Access to global device memory → go through 2 levels of cache
- ▶ L1\$ and shared memory corresponds to the same physical memory or distinct space (depending on version)
- ▶ *Read-only cache* (48 KB)



CUDA Programming

» Host-Side Programming

Overview

- ▶ Kernel-based programming
- ▶ Two parts:
 - Host programming
 - Compute kernels
- ▶ Need to transfer data back and forth
- ▶ Compilation chain dedicated to computer kernels (slightly different source language)
- ▶ Link to an software ecosystem
 - Runtime system (user space)
 - Driver (kernel space)

Programming Steps

- ▶ Template of CUDA code
 1. Initialization
 - Memory allocation
 - Read input data
 - ...
 2. GPU memory allocation
 - Dynamic allocation from host
 3. Host → device transfers
 4. Remote execution of compute kernel
 - Possibility to launch multiple kernels
 5. Device → host transfers

Compilation Process

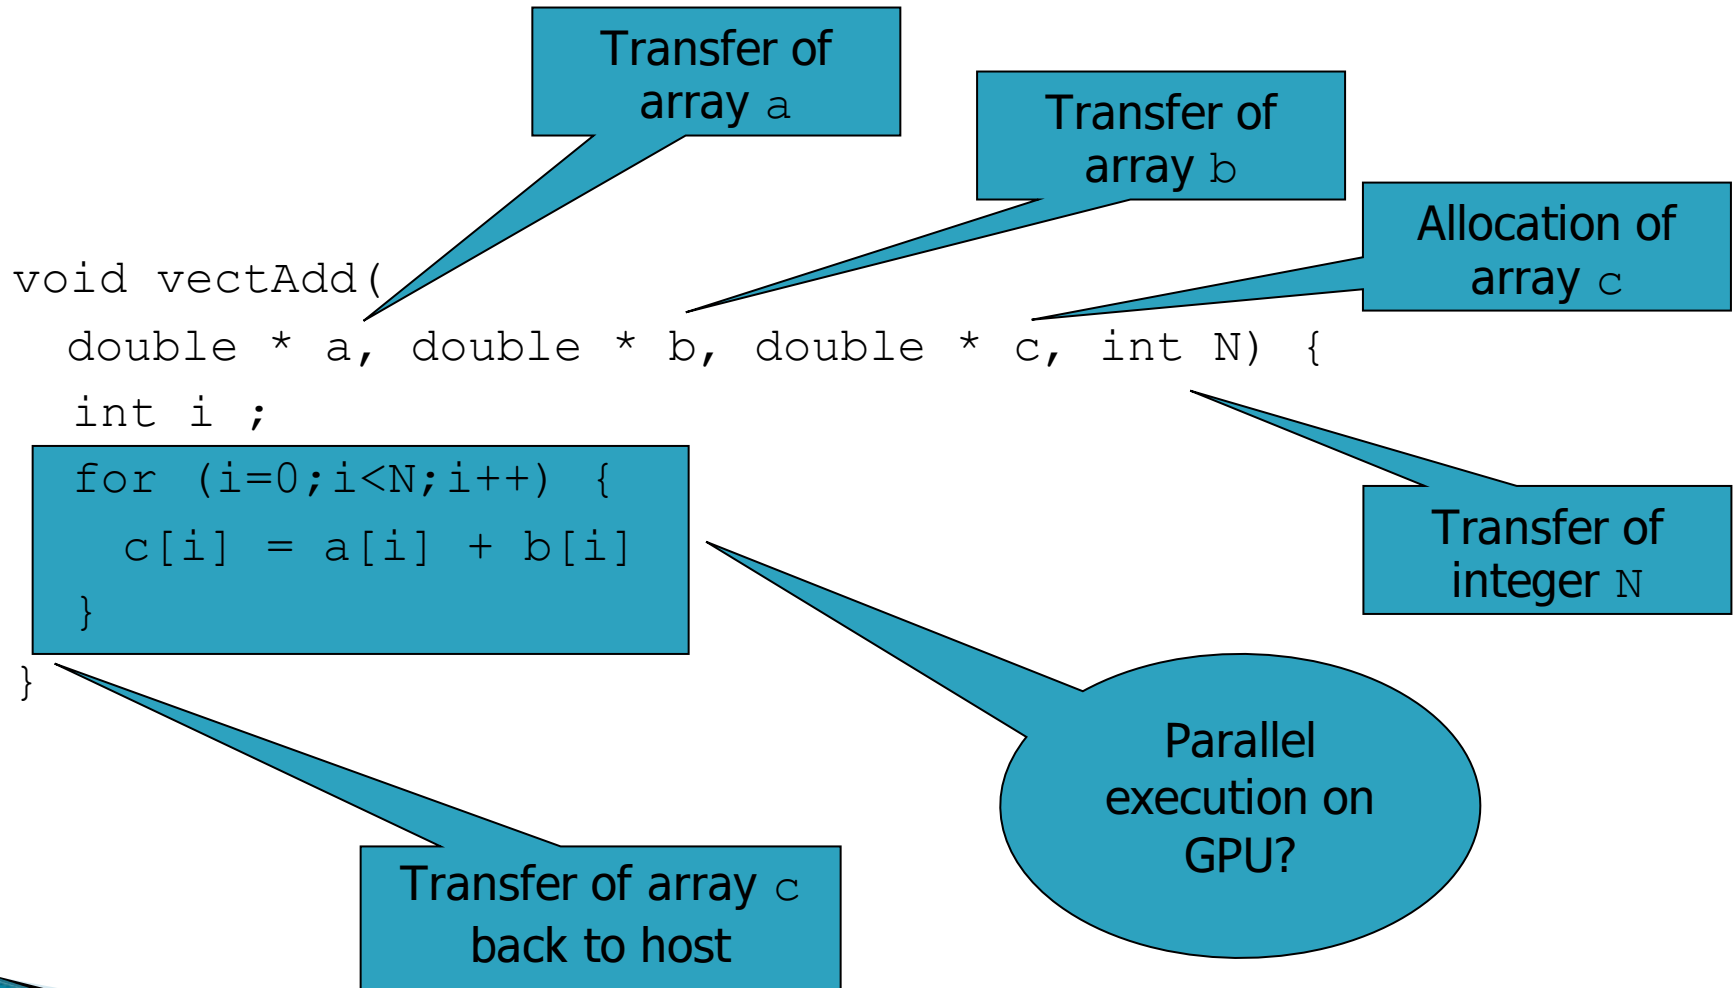
- ▶ Code structure
 - 2 parts
 - Host part
 - Device part
 - Two parts can be mixed inside same source file
- ▶ Host programming
 - C or C++
 - FORTRAN support developed by PGI
- ▶ Device programming
 - CUDA language
 - Superset of C99
- ▶ Compilation chain: NVCC
 - Use of `nvcc` compiler
 - Options available in the SDK documentation

First Example

- ▶ Example: vector addition
- ▶ Basic operation:
 - $c[i] = a[i] + b[i]$
- ▶ How to program this kernel in sequential?
 - Simple loop
 - Application of basic operation on each vector element

```
void vectAdd(  
    double * a, double * b,  
    double * c, int N) {  
    int i ;  
    for (i=0; i<N; i++) {  
        c[i] = a[i] + b[i] ;  
    }  
}
```


First Example



Host Programming

- ▶ Management from host (CPU)
 - Need to handle GPU global memory
 - Need to transfer data back and forth
 - Launch the remote device kernel
- ▶ Host-side CUDA API
 - CUDA high level / low level
 - *Runtime or driver*
 - `cudaMalloc`: allocate memory space on device
 - `cudaMemcpy`: transfer data with device
 - Multiple directions available

Host Programming

```
void vectAdd(  
    double * a, double * b,  
    double * c, int N) {  
  
    double * d_a ;  
    double * d_b ;  
    double * d_c ;  
  
    cudaMalloc((void **)&d_a,  
N*sizeof(double));  
    cudaMalloc((void **)&d_b,  
N*sizeof(double));  
    cudaMalloc((void **)&d_c,  
N*sizeof(double));
```

Device
pointers

Memory
allocation
on GPU

```
    cudaMemcpy(d_a, a,  
N*sizeof(double),  
    cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, b,  
N*sizeof(double),  
    cudaMemcpyHostToDevice);  
    cudaMemcpy(d_c, c,  
N*sizeof(double),  
    cudaMemcpyHostToDevice);
```

Transfe
r to
GPU

```
    vectAddKernel(d_a, d_b, d_c,  
N);
```

Kernel
execution

```
    cudaMemcpy(c, d_c,  
N*sizeof(double),  
    cudaMemcpyDeviceToHost);
```

```
/* cudaFree... */  
}
```

Transfer
to host

Execution Model

- ▶ Device only has to execute kernel: vector addition

- ▶ How to program kernel on device?

```
for (i=0; i<N; i++)  
    c[i]=a[i]+b[i];
```

- ▶ CUDA model defines notion of thread
 - Each thread will run the kernel function
 - By default, CUDA programs are multithreaded
 - Point of view from threads, and not from loops
 - Different from regular sequential programming

- ▶ Kernel pseudo-code:

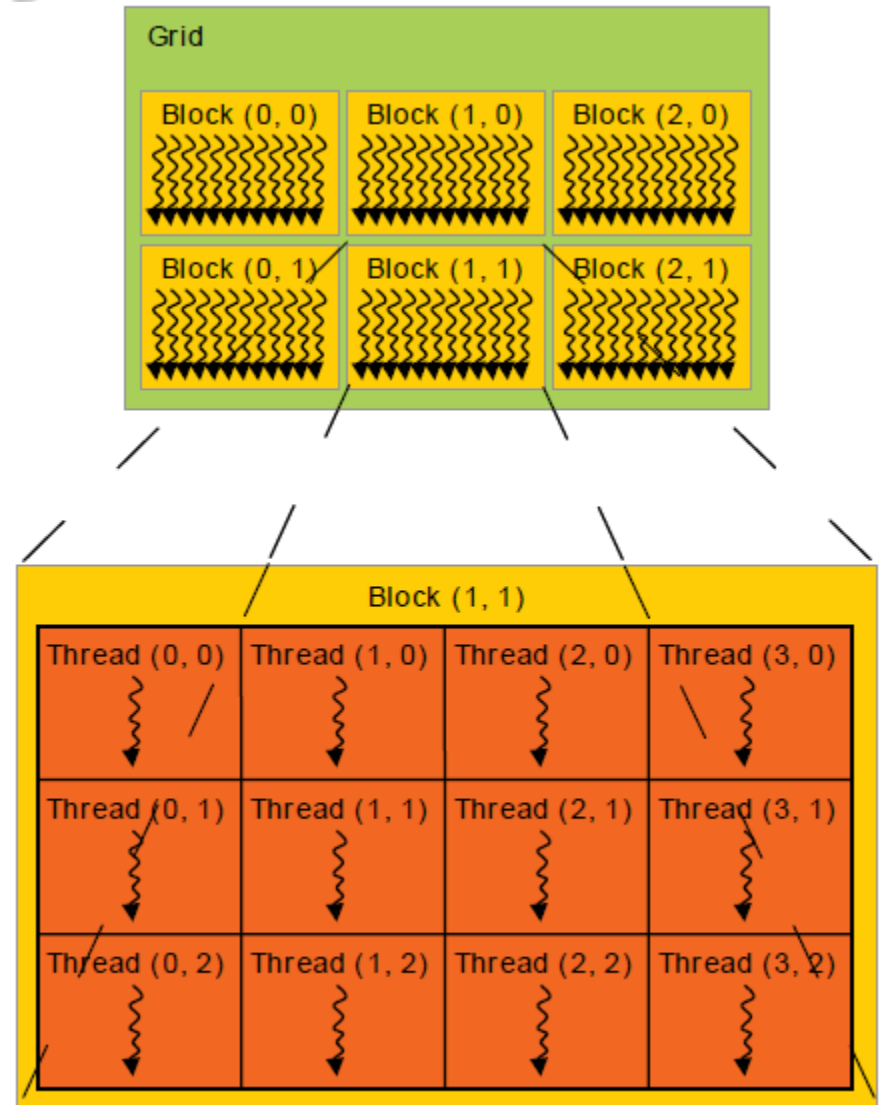
```
int i ;  
/* Update i with thread index */  
c[i] = a[i] + b[i];
```

Thread Hierarchy

- ▶ How to get thread id?
 - Related to thread organization w/ execution model
- ▶ CUDA introduces notion of blocks and grid to organize threads in hierarchical manner
- ▶ Grid
 - Contains blocks
 - Contains threads
- ▶ Grid and blocks may be organized in 1D, 2D or 3D
 - Logical representation

Thread Hierarchy

- ▶ Example
 - 1 2D grid
 - 6 2D blocks
- ▶ 6 blocks in the grid and 12 threads per block
 - ➔ total of 72 threads
- ▶ How to compute thread id from this representation?



CUDA Programming

- ▶ Kernel declaration
 - Keyword: `__global__`
- ▶ Compute thread index according to grid/block geometry
 - Example w/ 1D grid and 1D blocks
 - Organization should map the computational part
- ▶ Test to check that threads do not write beyond array size
 - Corresponds to upper bound from sequential loop

```
__global__ void
vecAddKernel( double *a,
double *b, double *c, int
N) {

    int i ;

    i = blockIdx.x *
blockDim.x + threadIdx.x
;

    if ( i < N ) {
        c[i] = a[i] + b[i];
    }
}
```


Kernel Integration

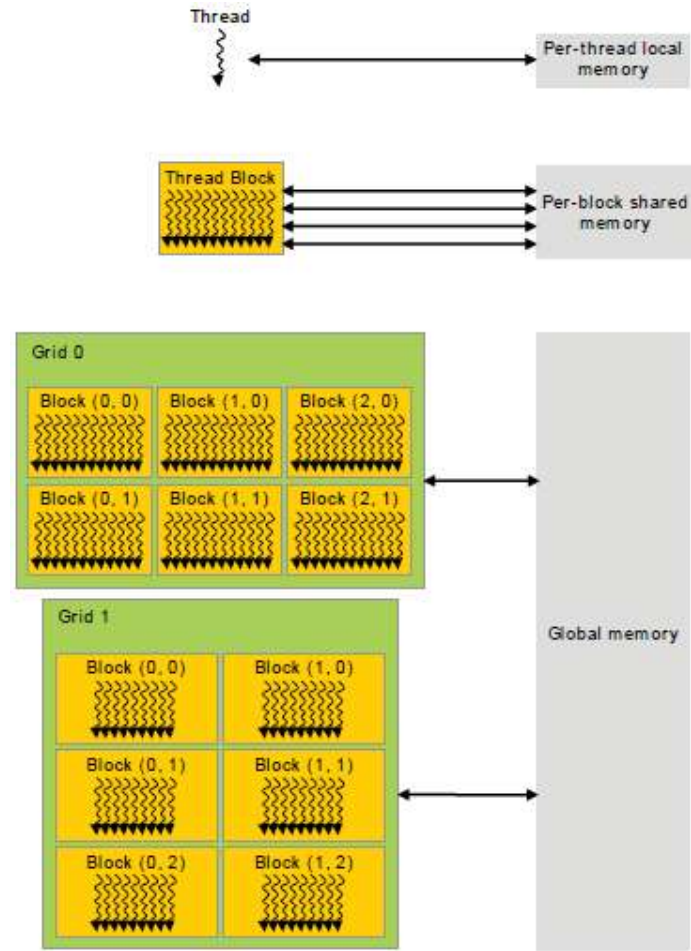
- ▶ How to choose representation of grid and blocks?
- ▶ Need to specify it during kernel remote launch
 - Host programming is in charge to define computational dimensions
- ▶ Syntax to launch a CUDA kernel from host to device

```
my_kernel<<<Dg, Db>>>(arg1, arg2, arg3 );
```

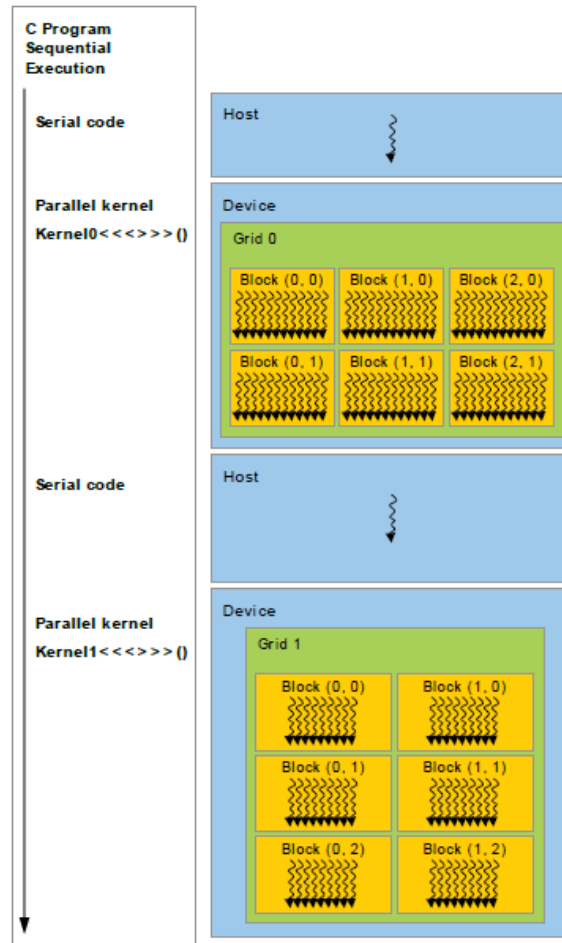
 - Dg: dimensions and size of grid (dim3 type)
 - Db: dimensions and size of blocks in grid (dim3 type)
- ▶ Total number of blocks?
- ▶ Number of thread per block?
- ▶ Total number of threads?

Memory Hierarchy

- ▶ Per-element vision
- ▶ Thread
 - Private local memory (stack)
 - Private register file
- ▶ Block
 - Private shared memory
- ▶ Grid
 - Global memory



Multi-Kernel Execution Model



Asynchronism

- ▶ By default some functions have asynchronous behavior
 - Kernel execution
 - Device-to-device copies
 - Memory initialization
- ▶ Possibility to wait for the end of operation execution
 - Sort of barrier

```
cudaThreadSynchronize();
```
- ▶ Force synchronism
 - `CUDA_LAUNCH_BLOCKING=1`
- ▶ Consecutive calls to kernel *vecAdd*
 - On different vectors?
 - If RAW dependencies exists?

Advanced Asynchronism

▶ Motivations

- Ability to transfer data during kernel execution
- Ability to execute multiple compute kernels at the same time
 - If target device allows it (starting w/ Fermi)
 - If there are no data dependences between simultaneous kernels

▶ Solution: *streaming*

▶ Stream declaration

- Interaction w/ device on a particular stream
- Driver is aware of what can be done in parallel
- `cudaMemcpy(d_c, c, N*sizeof(double), cudaMemcpyHostToDevice, stream[0]);`
- `my_kernel<<<Dg, Db, stream[0]>>>(arg1, arg2, arg3);`

Profiling and Timers

- ▶ Profiling
 - Use provided tools like `nvprof` (see Lab)
 - Use CUDA events to track timing in CUDA functions
- ▶ Main type: `cudaEvent_t`
- ▶ Event creation
 - `cudaEventCreate(cudaEvent_t * e)`
- ▶ Event activation
 - `cudaEventRecord(cudaEvent_t e, cudaStream_t s)`
- ▶ Waiting for even to be activated
 - `cudaEventSynchronize(cudaEvent_t e);`
- ▶ Compute time spent between events
 - `cudaEventElapsedTime(float * ms, cudaEvent_t start, cudaEvent_t stop);`

Error Handling

- ▶ In CUDA, almost all functions return an error code
 - Return type: `cudaError_t`
- ▶ If everything is ok → `cudaSuccess`
- ▶ Otherwise, possibility to access to error description
 - `const char * cudaGetErrorString (cudaError_t error);`
- ▶ For functions not returning error information (e.g., kernel call)
 - Access to last error through `cudaGetLastError()`
 - Return type: `cudaError_t`

Debugging

- ▶ How to debug a CUDA code?
- ▶ Add calls to `printf`?
 - Possible but not fully safe
- ▶ Add synchronization
- ▶ Check all error codes
 - Main way to check errors
- ▶ Transfer back to host on regular basis
 - Allow full checking w/ printing...

CUDA Programming

»» Device-Side Programming

CUDA Programming

- ▶ Compute kernel
 - Based on C99 standard
 - Some restrictions
 - Some extensions
- ▶ Extensions
 - New keywords
 - Variables defined by default
 - New function
- ▶ See *CUDA C Programming Guide* document
 - Appendix B for language extension
 - Appendix F for C/C++ standard support

Basic Kernel

- ▶ Basic syntax for a compute kernel
 - Function with *void* return type
 - Attribute to declare the function to execute on device
 - `__global__`
 - Input arguments

- ▶ Example

```
__global__ void vecAddKernel( double *a,  
    double *b, double *c, int N ) {  
    int i ;  
    i = blockIdx.x * blockDim.x + threadIdx.x ;  
    if ( i < N ) {  
        c[i] = a[i] + b[i];  
    }  
}
```

Compilation

- ▶ Use dedicated compiler chain: NVCC
- ▶ Compiler will be in charge of splitting code parts for the host and device
- ▶ File extensions
 - .cu → file containing CUDA code (and host code)
 - .c / .cpp → file containing only host code (including CUDA API)

Extensions – Keywords

- ▶ Definition of new keywords
- ▶ Categories
 - Function attributes
 - Variable attributes
 - Types
- ▶ Set of variables defined by default
 - Rely on new defined data types

Function Attributes

- ▶ Keyword to add during function declaration and definition
 - Coherency between multiple declarations and definition!
 - Between return type and function name
- ▶ Function executed on device and callable from host (i.e., starting point for kernels)
 - `__global__`
- ▶ Function dedicated to host or device (combinable)
 - `__host__`
 - `__device__`
- ▶ Behavior by default → equivalent to `__host__`

Function Restrictions

- ▶ Functions with `__global__` attribute
 - Return type `void`
 - Call with execution context (number/dimensions of blocks, number of threads per block...)
 - Asynchronous call
 - Arguments stores in shared memory (256B) or constant memory (4KB)
 - Depends on CUDA capabilities of the target device
 - Impossible to get its address (function pointer)
- ▶ Functions executing on device
 - No static variable
 - No variable number of arguments
 - Restricted recursion (only for functions declared with `__device__` attribute)

Variable Attribute

- ▶ New variable attributes to handle placement
- ▶ Device-resident variable
 - `__device__`
 - By default in global memory (shared by all threads during application lifetime)
- ▶ Constant memory resident variable
 - `__constant__`
 - Lifetime of application
 - Cannot be defined on device
- ▶ Shared-memory variable
 - `__shared__`
 - Shared between threads in the same block
 - One copy per block
 - Lifetime of block
- ▶ By default, variable declared on device is stored in a register

Volatile Variables

- ▶ Shared-data synchronization

- ▶ Example of concurrent accesses

```
// myArray is an array of non-zero integers
// located in global or shared memory
__global__ void MyKernel(int* result) {
    int tid = threadIdx.x;
    int ref1 = myArray[tid] * 1;
    myArray[tid + 1] = 2;
    int ref2 = myArray[tid] * 1;
    result[tid] = ref1 * ref2;
}
```

- ▶ What contains `result[tid]` ?
 - `myArray[tid]` is stored in a register → `ref1==ref2`
- ▶ If declared volatile → ok (or put a *memory fence*)
 - No guarantee on execution order

Variable Restrictions

- ▶ Dynamic management of variables in shared memory

```
extern __shared__ char array[];
__device__ void func() {
    short* array0 = (short*)array;
    float* array1 = (float*)&array[128];
    int* array2 = (int*)&array[64];
}
```

- ▶ Need to handle by hand placement of data in shared memory
 - No dynamic allocation
 - Need to respect alignment rules

Data Types

▶ New types

- Vectors
- Multi-dimension integers

▶ Vectors

- Name convention: base type + width
- Examples: `int2`, `float4`
- Need to respect alignment rules
- Functions available to create such types
 - Example: `int2 make_int2(int x, int y);`

Data Types

▶ 3D integers

- `dim3`
- Equivalent to vector type `uint3`
- Access to coordinates through fields `x`, `y` and `z`
- By default, initialized to 1

▶ Example

```
dim3 a ;  
a.x = 4 ;
```

Predefined Variables

- ▶ **Grid dimensions**
 - `dim3 gridDim`
 - Contains number of blocks along 3 dimensions
- ▶ **Index of block in grid**
 - `dim3 blockIdx`
 - Contains coordinates of current block in grid
- ▶ **Block dimensions**
 - `dim3 blockDim`
 - Contains number of threads along 3 dimensions
- ▶ **Index of thread in block**
 - `uint3 threadIdx`
 - Contains coordinates of current thread in its blocks
- ▶ **Warp size**
 - `int warpSize`

Available Functions

- ▶ Memory barrier
- ▶ Synchronizations
- ▶ Math operations
- ▶ Atomic operations
- ▶ *Timing*
- ▶ I/O

Memory Barrier

- ▶ Weakly-ordered memory model
 - Order of store instructions may be viewed as different from one thread to another
- ▶ Memory barrier
 - Mechanism to guarantee that memory accesses performed before the barrier appear to all threads as launched before the barrier
- ▶ Different levels of barrier
 - Block: `void __threadfence_block();`
 - Device: `void __threadfence();`
 - Device + host: `void __threadfence_system();`

Synchronization

- ▶ Synchronization among all threads in the same block
 - `void __syncthreads();`
 - Include data synchronization
 - Be careful of control flow !
 - Act as a collective operations within a block
- ▶ For devices w/ CUDA capabilities ≥ 2.0
 - `int __syncthreads_count(int predicate);`
 - `int __syncthreads_and(int predicate);`
 - `int __syncthreads_or(int predicate);`

Math Operations

- ▶ Set of math functions optimized for GPU
 - Ex: `sin(float)`, `cos(double)`, ...
 - Less instructions but less accurate!
- ▶ Functions with `__` prefix
 - `sinf(x) → __sinf(x)`
- ▶ Compilation option to automatically use these functions
 - Transform regular calls to optimized ones
 - `-use_fast_math`
 - Mainly available for single-precision computations

Atomic Operations

- ▶ Operation in one instruction that guarantee atomicity
 - E.g., `int atomicAdd(int* address, int val);`
- ▶ Available functions
 - Addition: `atomicAdd`
 - Substraction: `atomicSub`
 - Exchange: `atomicExch`
 - Min/Max: `atomicMin`
 - Increment/Decrement: `atomicInc`
 - CAS, ...
- ▶ Restrictions
 - Available for integers
 - Available for single-precision floats (capabilities 2.X)
 - Since 6.X, double-precision operations

Timing & I/O

▶ Multiprocessor timestamp

- `clock_t clock();`
- Number of cycles
- Possibility to check number of clocks for a thread

▶ Formatted output

- `int printf(const char *format[, arg, ...]);`
- Require at least CUDA capabilities 2.X
- Per-thread context
- Final formatting on the host

Example

- ▶ Matrix multiply
- ▶ Rely on C data structure
 - Available on CPU and GPU
- ▶ Use variables w/ 2 dimensions
- ▶ Every thread computes one element of resulting matrix

```
typedef struct {  
    int width;  
    int height;  
    float* elements;  
} Matrix;
```

- ▶ What if matrices are larger than maximum number of allowed threads?

```
__global__ void MatMulKernel(Matrix  
    A, Matrix B, Matrix C)  
{  
  
    float Cvalue = 0;  
    int row = blockIdx.y * blockDim.y +  
        threadIdx.y;  
    int col = blockIdx.x * blockDim.x +  
        threadIdx.x;  
  
    for (int e = 0; e < A.width; ++e)  
        Cvalue += A.elements[row *  
            A.width + e]  
            * B.elements[e * B.width +  
                col];  
    C.elements[row * C.width + col] =  
        Cvalue;  
}
```

Example

```
__global__ void MatMulKernel(Matrix A, Matrix B, Matrix C)
{

    float Cvalue = 0;
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;

    for (int r = row ; r < C.height ; r += blockDim.y )
        for (int c = col ; c < C.width ; c += blockDim.x )
            for (int e = 0; e < A.width; ++e)
                Cvalue += A.elements[row * A.width + e]
                    * B.elements[e * B.width + col];
            C.elements[row * C.width + col] = Cvalue;

}
```

Register Allocation

- ▶ Each compute kernel uses several registers
- ▶ Depending on instructions inside the kernel body
 - Compiler transformations/optimizations
 - Register allocation
- ▶ But
 - Limited size of register file
 - Register file is shared among threads inside the same *Streaming Multiprocessor*
- ▶ Relation w/ number of threads?

Register Allocation

- ▶ Option to set an upper bound to compiler
 - `-maxrregcount=N`
- ▶ Attribute to indicate max number of threads and blocks when calling target function

```
__global__ void  
__launch_bounds__(maxThreadsPerBlock,  
minBlocksPerMultiprocessor)  
MyKernel(...)  
{  
    ...  
}
```

Memory Access

- ▶ Because of synchronous execution
 - Thread executing load operation may lower other's speed
 - Several memory accesses may be serialized
 - Long-latency access → need more warps to recover latency
- ▶ Possible optimizations
 - *Load coalescing*
 - Avoid bank conflicts

Load coalescing

- ▶ Access to global memory
- ▶ Concurrent accesses emitted by threads inside the same warp
 - All threads inside the same warp execute the same instruction at the same time
- ▶ Memory requests are serialized w/ chunks of 128B
 - Size of cache line
 - Optimization is memory accesses are contiguous

Load coalescing and shared memory

- ▶ Shared memory
 - Need to explicitly declare variables inside shared memory
 - Data transfers at the beginning of kernel
 - Update to global memory at the end of kernel
- ▶ Optimization: benefit from memory-to-memory transfers to change data layout
 - Even if all data are not necessary
- ▶ Be careful to bank conflicts

Best Practices

- ▶ Higher priority
 - Think parallel
 - Minimize data transfers between host and device
 - Access main device memory w/ coalesced operations
 - Use shared memory when possible
 - Avoid path divergence
- ▶ Lower priority
 - Avoid bank conflict when accessing shared memory
 - Avoid too large number of threads per blocks (multiple of 32)
 - Rely on optimized math functions